

Big Data Analytics

Map-Reduce and Spark

Slides by:

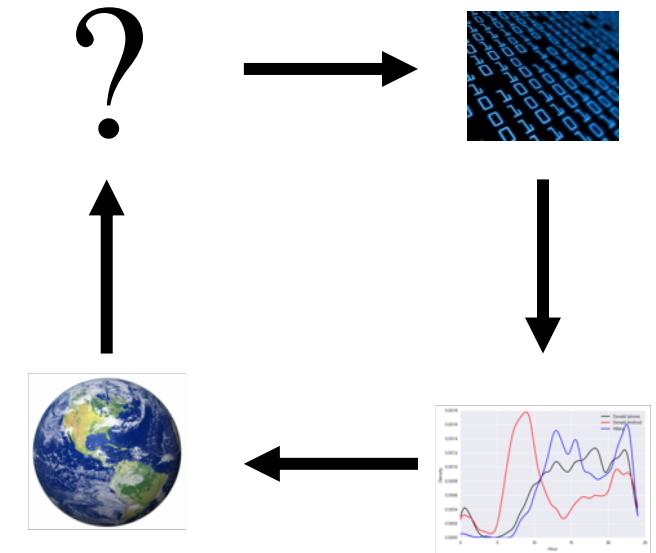
Joseph E. Gonzalez

jegonzal@cs.berkeley.edu

With revisions by:

Josh Hug

hug@cs.Berkeley.edu



From SQL to Big Data (with SQL)

- Last week...
 - Databases
 - (Relational) Database Management Systems
 - SQL: Structured Query Language
- Today
 - More on databases and database design
 - Enterprise data management and the data lake
 - Introduction to distributed data storage and processing
 - Spark

Data in the Organization

A little bit of buzzword bingo!

OLAP (Online Analytics Processing)

Operational Data Store

Star Schema

Data Lake

CUBE

Data Warehouse

ROLLUP

Drill Down

ETL (Extract, Transform, Load)

Snowflake Schema

Schema on Read

Inventory



How we like to think of data in the organization

The reality...



Sales
(Asia)



Inventory



Sales
(US)



Advertising





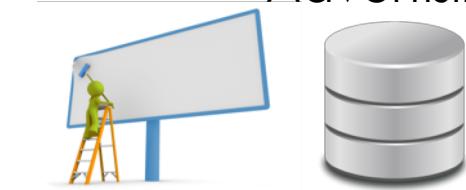
Sales
(Asia)



Sales
(US)



Inventory

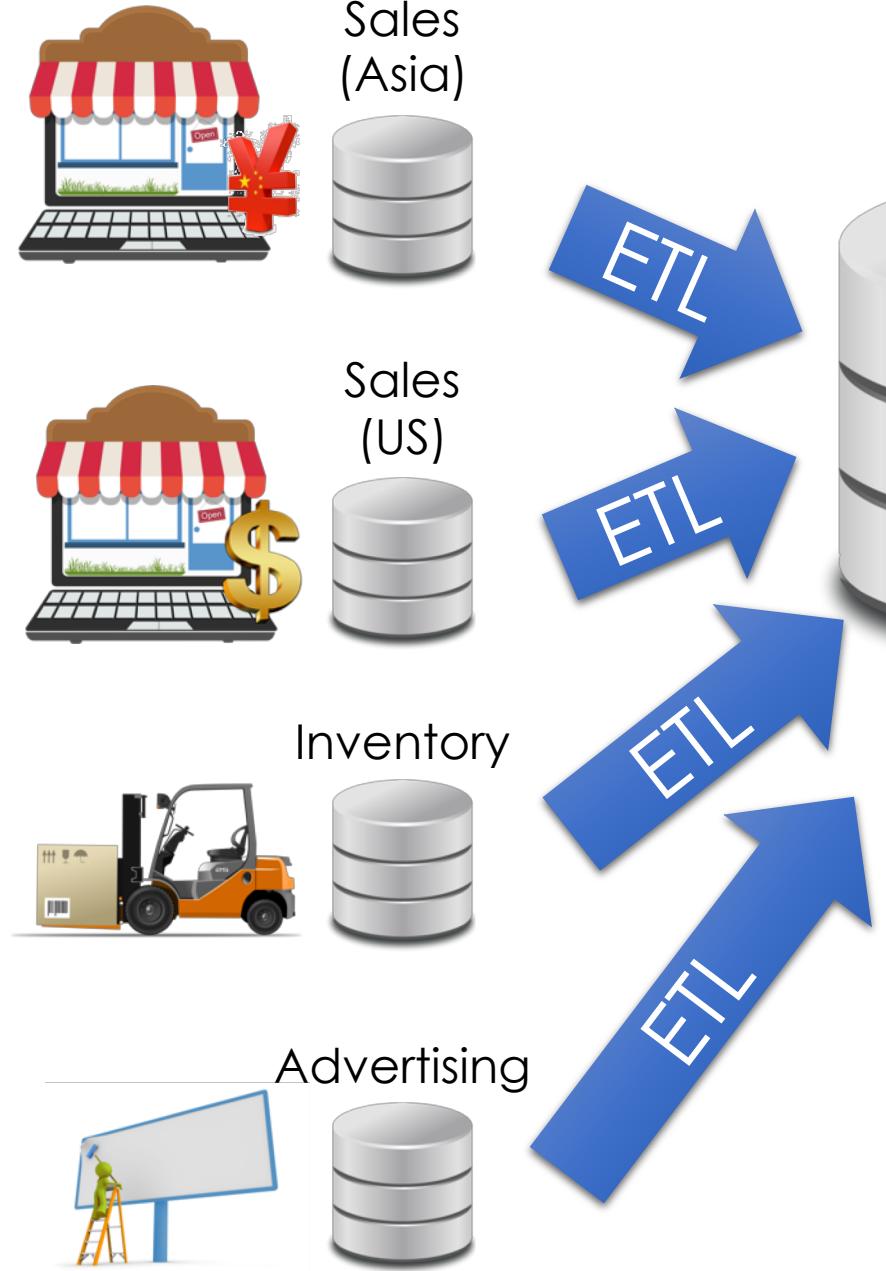


Advertising

Operational Data Stores

- Capture **the now**
- Many different databases across an organization
- Mission critical... be careful!
 - Serving live ongoing business operations
 - Managing inventory
- Different formats (e.g., currency)
 - Different schemas (acquisitions ...)
- Live systems often don't maintain history

We would like a consolidated, clean, historical snapshot of the data.



Data Warehouse

Collects and organizes historical data from multiple sources

Data is periodically **ETL**ed into the data warehouse:

- **Extracted** from remote sources
- **Transformed** to standard schemas
- **Loaded** into the (typically) relational (SQL) data system

Extract → Transform → Load (ETL)

Extract & Load: provides a snapshot of operational data

- Historical snapshot
- Data in a single system
- Isolates analytics queries (e.g., Deep Learning) from business critical services (e.g., processing user purchase)
- Easy!

Transform: clean and prepare data for analytics in a unified representation

- **Difficult** → often requires specialized code and tools
- Different schemas, encodings, granularities



Data Warehouse

Collects and organizes historical data from multiple sources

How is data organized in the Data Warehouse?

Example Sales Data

pname	category	price	qty	date	day	city	state	country
Corn	Food	25	25	3/30/16	Wed.	Omaha	NE	USA
Corn	Food	25	8	3/31/16	Thu.	Omaha	NE	USA
Corn	Food	25	15	4/1/16	Fri.	Omaha	NE	USA
Galaxy	Phones	18	30	1/30/16	Wed.	Omaha	NE	USA
Galaxy	Phones	18	20	2/28/16	Thu.	Omaha	NE	USA
Galaxy	Phones	18	50	3/31/16	Fri.	Omaha	NE	USA
Peanuts	Food	2	45	3/31/16	Wed.	Omaha	NE	USA
Peanuts	Food	2	45	4/1/16	Thu.	Seoul		Korea

Multidimensional Data Model

Sales **Fact Table**

pid	timeid	locid	sales
11	1	1	25
11	2	1	8
11	3	1	15
12	1	1	30
12	2	1	20
12	3	1	50
12	1	1	8
13	2	1	10
13	3	1	10
11	1	2	35
11	2	2	22
11	3	2	10
12	1	2	26

Locations

locid	city	state	country
1	Omaha	Nebraska	USA
2	Seoul		Korea
5	Richmond	Virginia	USA

Products

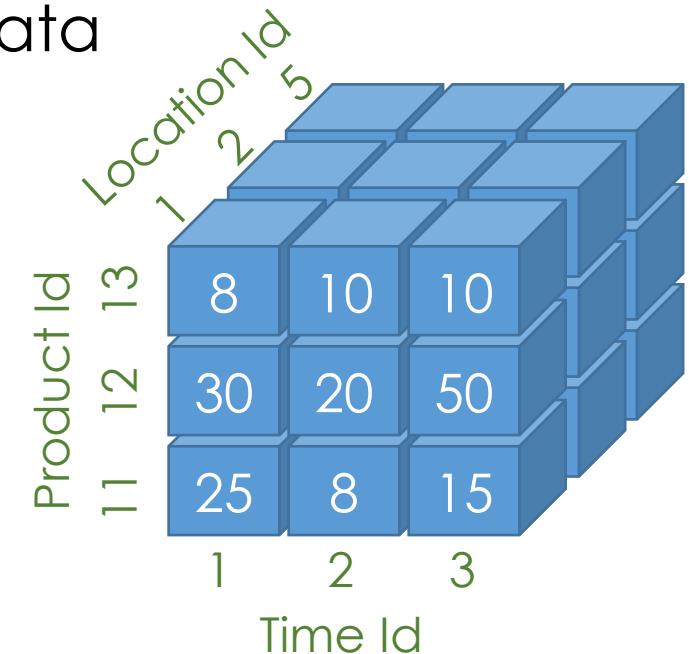
pid	pname	category	price
11	Corn	Food	25
12	Galaxy 1	Phones	18
13	Peanuts	Food	2

Time

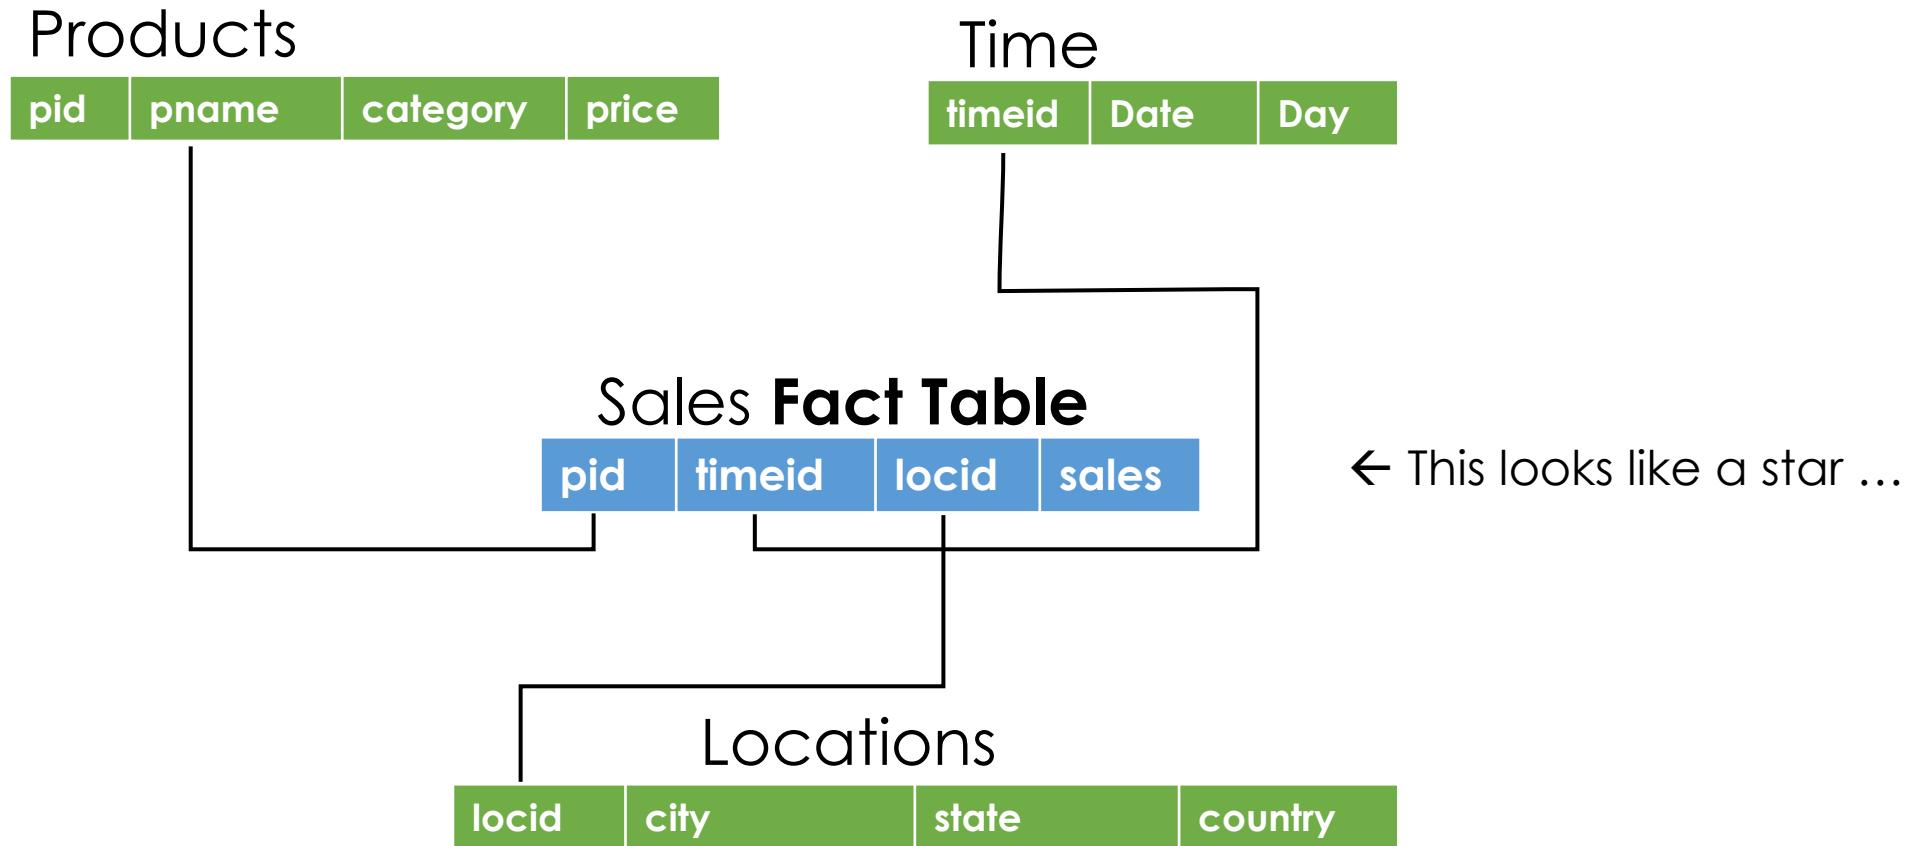
timeid	Date	Day
1	3/30/16	Wed.
2	3/31/16	Thu.
3	4/1/16	Fri.

Dimension Tables

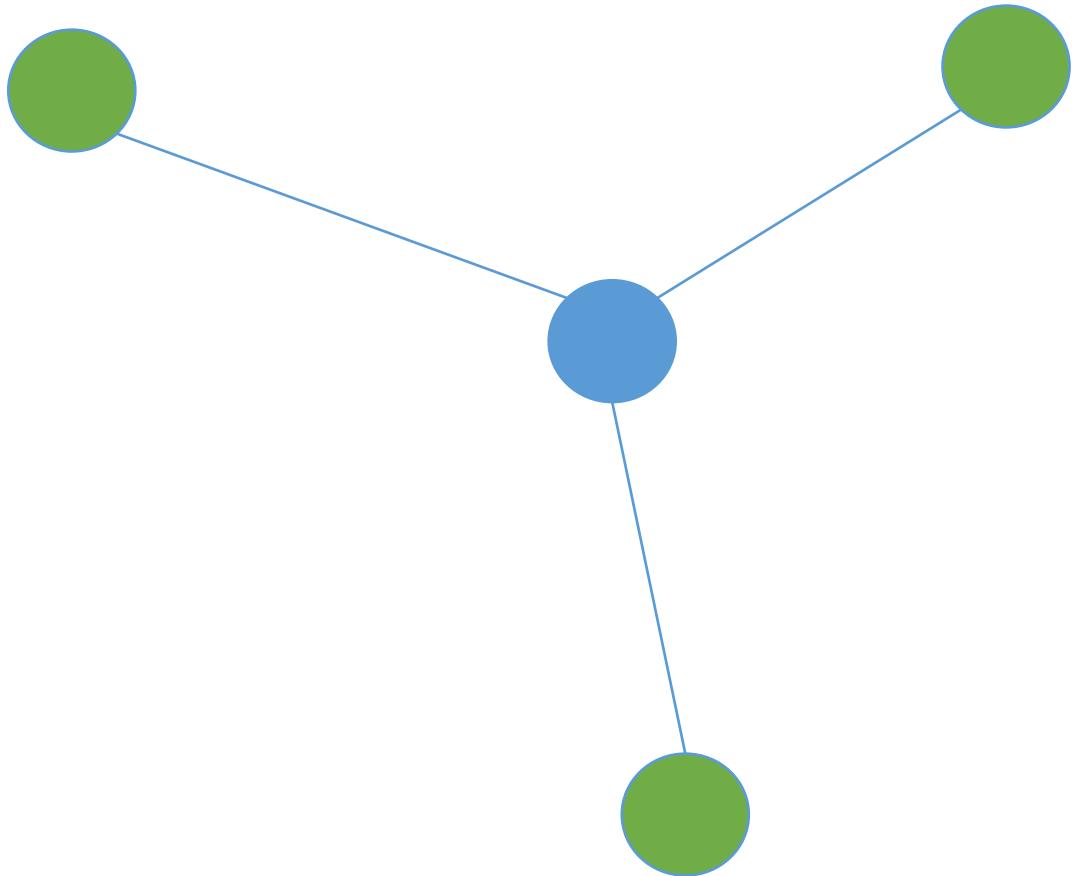
➤ Multidimensional “Cube” of data



The Star Schema

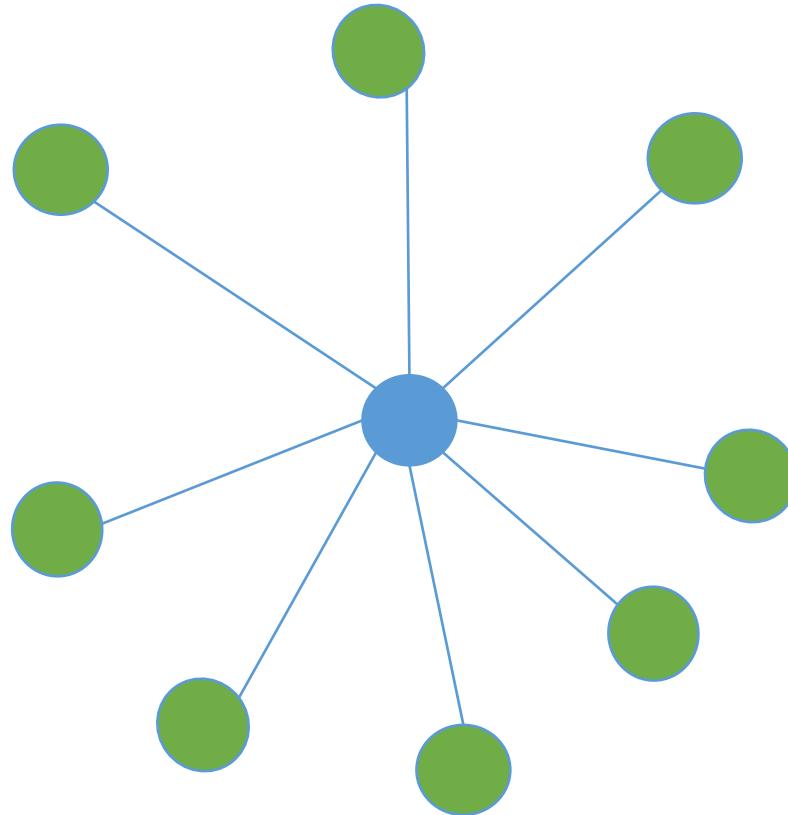


The Star Schema



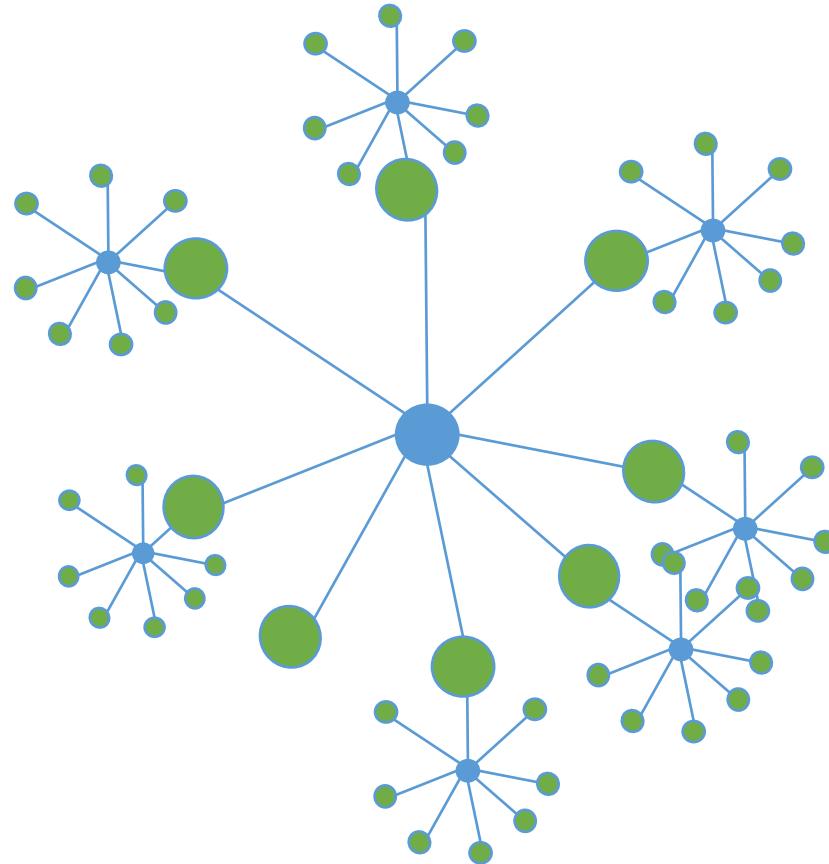
← This looks like a star ...

The Star Schema



← This looks like a star ...

The Snowflake Schema



← This looks like a snowflake ...

See CS 186 for more!

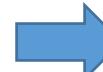
Online Analytics Processing (**OLAP**)

Users interact with multidimensional data:

- Constructing ad-hoc and often complex SQL queries
- Using graphical tools that to construct queries
 - e.g. Tableau

Let's discuss some common types of queries used in OLAP.

Cross Tabulation (Pivot Tables)



Item	Color	Quantity
Desk	Blue	2
Desk	Red	3
Sofa	Blue	4
Sofa	Red	5

		Item		
		Desk	Sofa	Sum
Color	Blue	2	4	6
	Red	3	5	8
	Sum	5	9	14

- Aggregate data across pairs of dimensions
 - **Pivot Tables:** graphical interface to select dimensions and aggregation function (e.g., SUM, MAX, MEAN)
 - **GROUP BY** queries
- Related to contingency tables and marginalization in stats.
- What about many dimensions?

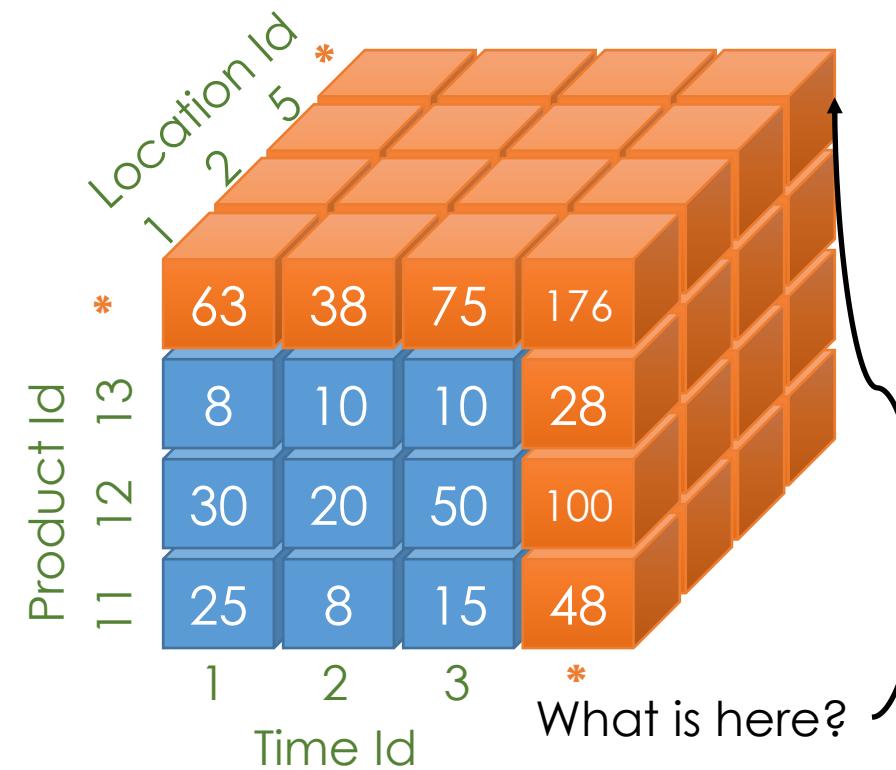
Cube Operator

- Generalizes cross-tabulation to higher dimensions.

- In SQL:

```
SELECT Item, Color, SUM(Quantity) AS QtySum  
FROM Furniture  
GROUP BY CUBE (Item, Color);
```

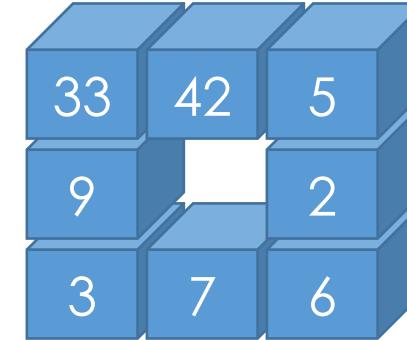
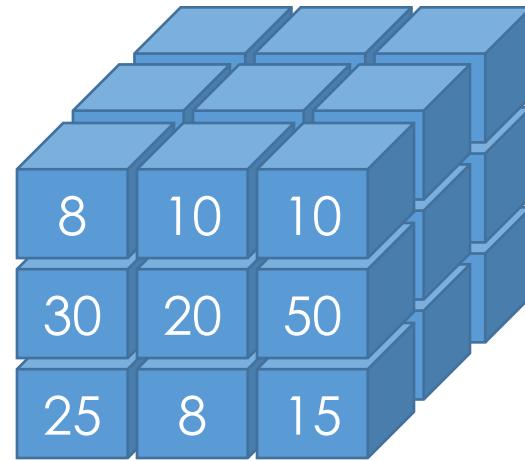
Item	Color	Quantity
Desk	Blue	2
Desk	Red	3
Sofa	Blue	4
Sofa	Red	5



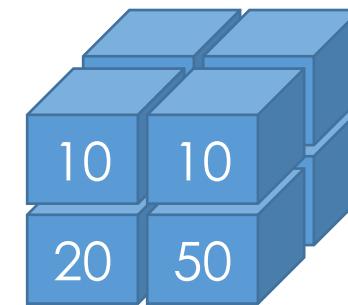
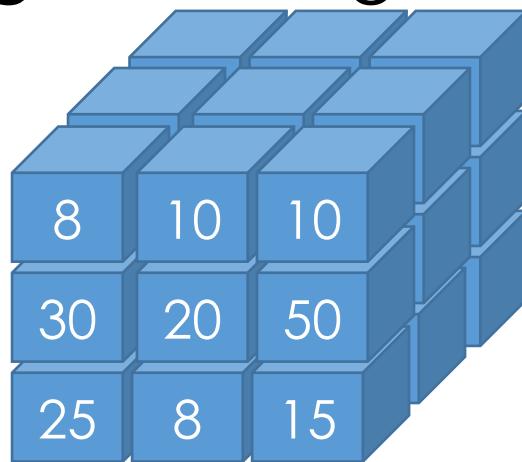
Item	Color	QtySum
Desk	Blue	2
Desk	Red	3
Desk	*	5
Sofa	Blue	4
Sofa	Red	5
Sofa	*	9
*	*	14
*	Blue	6
*	Red	8

OLAP Queries

- **Slicing:** selecting a value for a dimension

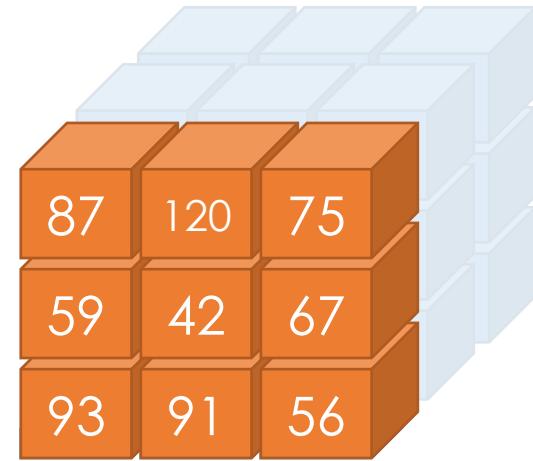
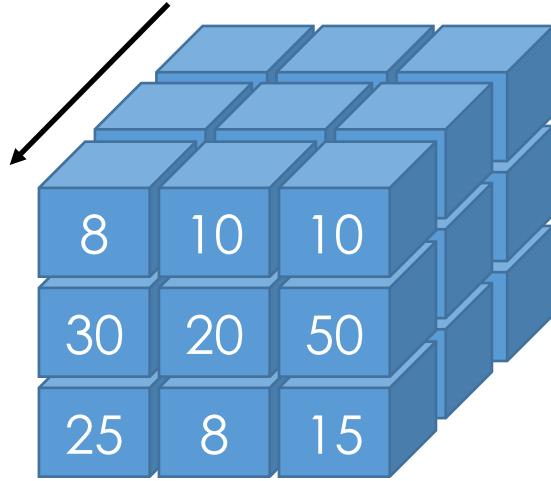


- **Dicing:** selecting a range of values in multiple dimension



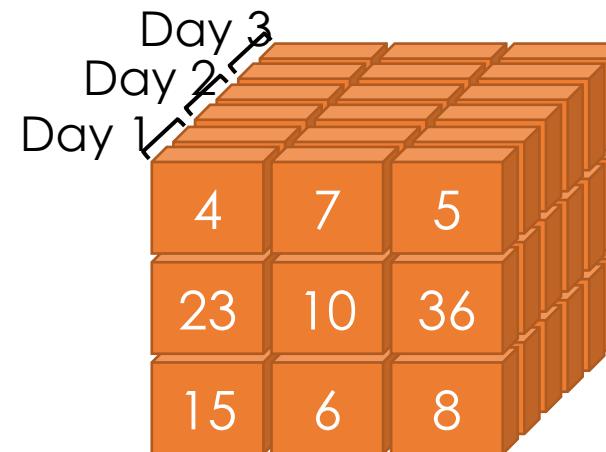
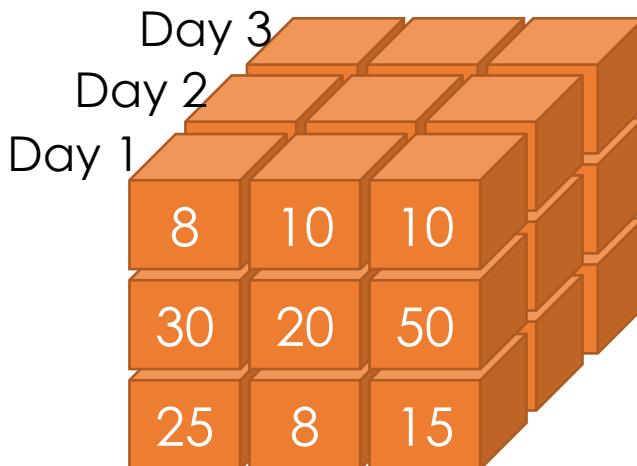
OLAP Queries

- **Rollup:** Aggregating along a dimension



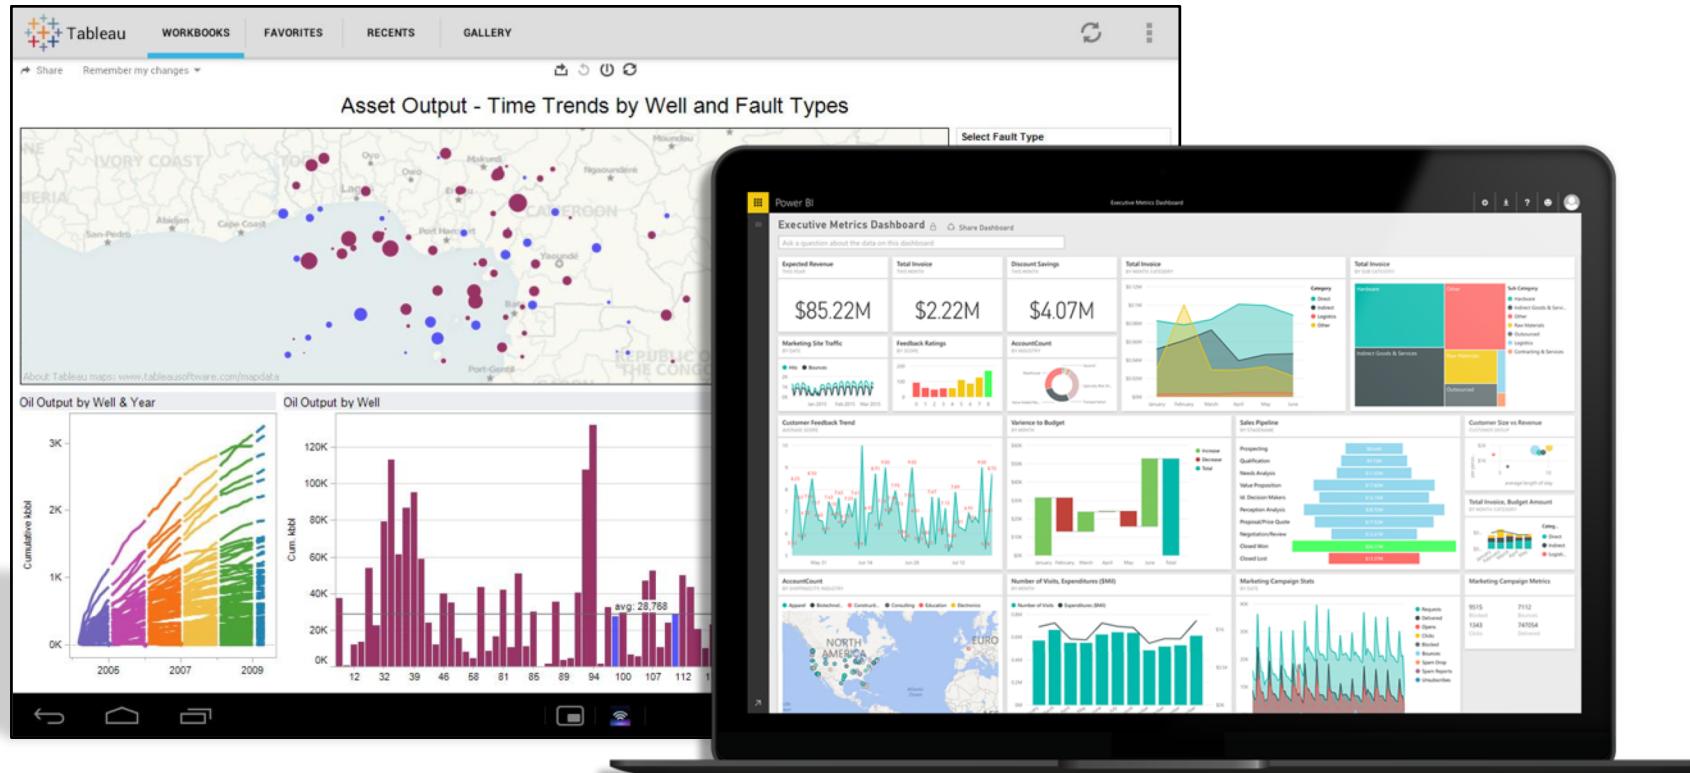
Similar to CUBE

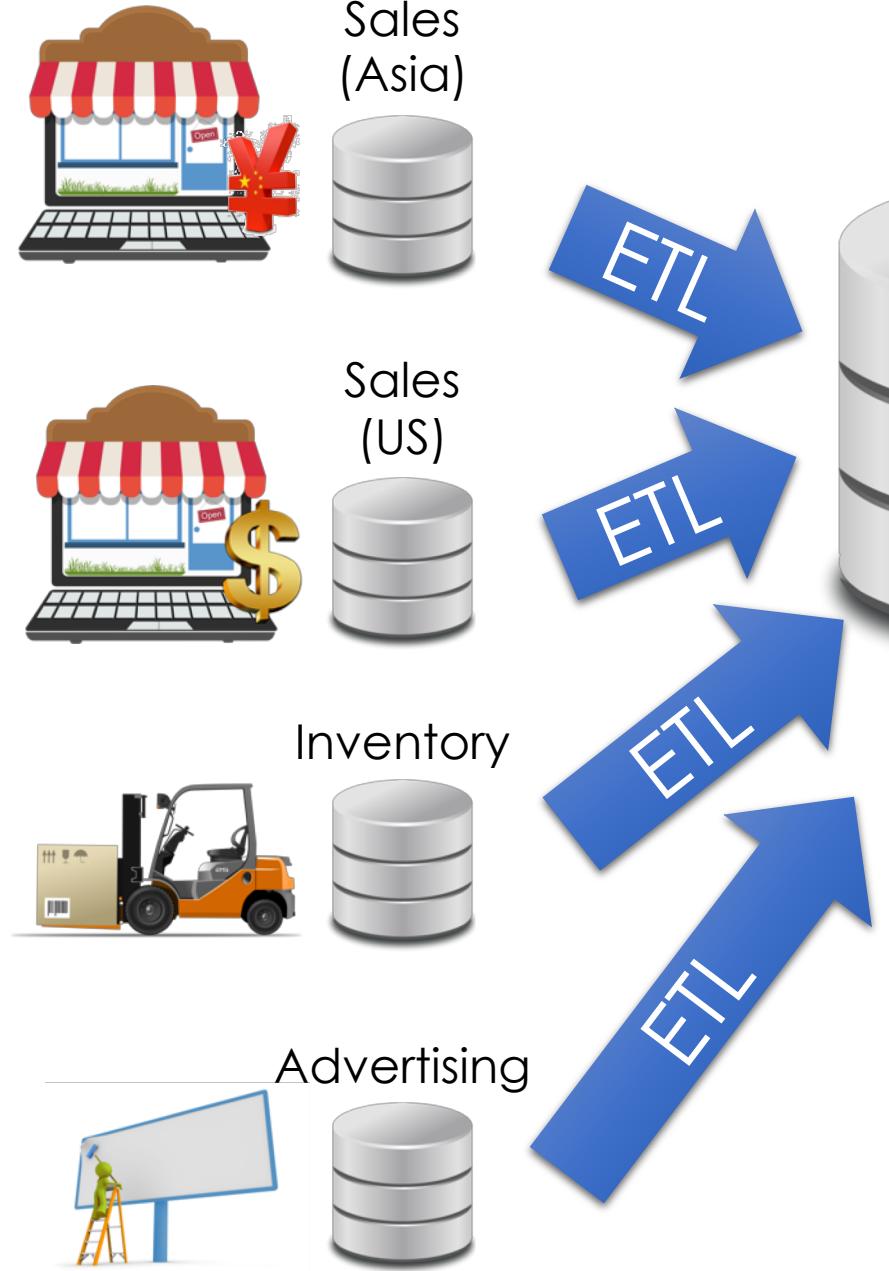
- **Drill-Down:** de-aggregating along a dimension



Reporting and Business Intelligence (BI)

- Use high-level tools to interact with their data:
 - Automatically generate SQL queries
 - Queries can get big!
- Common!





Data Warehouse

Collects and organizes historical data from multiple sources

So far ...

- Star Schemas
- Data cubes
- OLAP Queries



Data Warehouse

Collects and organizes historical data from multiple sources

- How do we deal with semi-structured and unstructured data?
- Do we really want to force a schema on load?

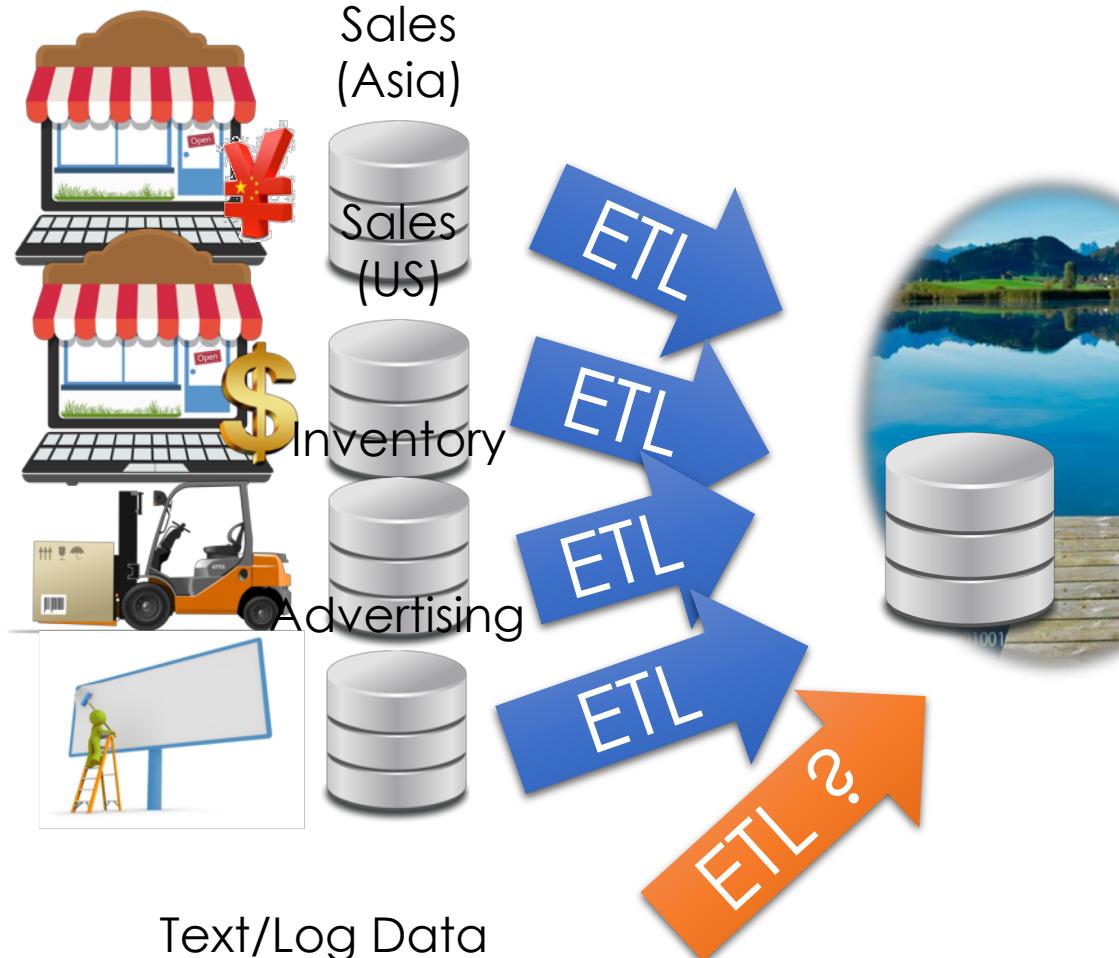
Data Warehouse

iid	date_taken	is_cat	is_grumpy	image_data
45123 1333	01-22-2016	1	1	
47234 2122	06-17-2017	0	1	
57182 7231	03-15-2009	0	0	
23847 2733	05-18-2018	0	0	

Unclear what a good schema for this image data might look like. Something like above will work, but it is inflexible!

Collects and organizes historical data from multiple sources

- How do we deal with semi-structured and unstructured data?
- Do we really want to force a schema on load?



Data Lake*

Store a copy of all the data

- in one place
- in its original “natural” form

Enable data consumers to choose how to transform and use data.

- Schema on Read

What could go wrong?

*Still being defined...[Buzzword Disclaimer]

The Dark Side of Data Lakes



- Cultural shift: *Curate* → *Save Everything!*
 - Noise begins to dominate signal
- Limited data governance and planning
 - **Example:** `hdfs://important/joseph_big_file3.csv_with_json`
 - **What** does it contain?
 - **When** and **who** created it?
- No cleaning and verification → lots of dirty data
- New tools are more complex and old tools no longer work

Enter the data scientist

A Brighter Future for Data Lakes

Enter the data scientist

- Data scientists bring new skills
 - Distributed data processing and cleaning
 - Machine learning, computer vision, and statistical sampling
- Technologies are improving
 - SQL over large files
 - Self describing file formats (e.g. Parquet) & catalog managers
- Organizations are evolving
 - Tracking data usage and file permissions
 - New job title: data engineers



How do we **store** and **compute** on large unstructured datasets

- Requirements:
 - Handle very **large files** spanning **multiple computers**
 - Use **cheap** commodity devices that **fail frequently**
 - **Distributed data processing** quickly and **easily**
- Solutions:
 - **Distributed file systems** → spread data over multiple machines
 - Assume machine **failure** is common → **redundancy**
 - **Distributed computing** → load and process files on multiple machines concurrently
 - Assume machine **failure** is common → **redundancy**
 - **Functional programming** computational pattern → **parallelism**

Distributed File Systems

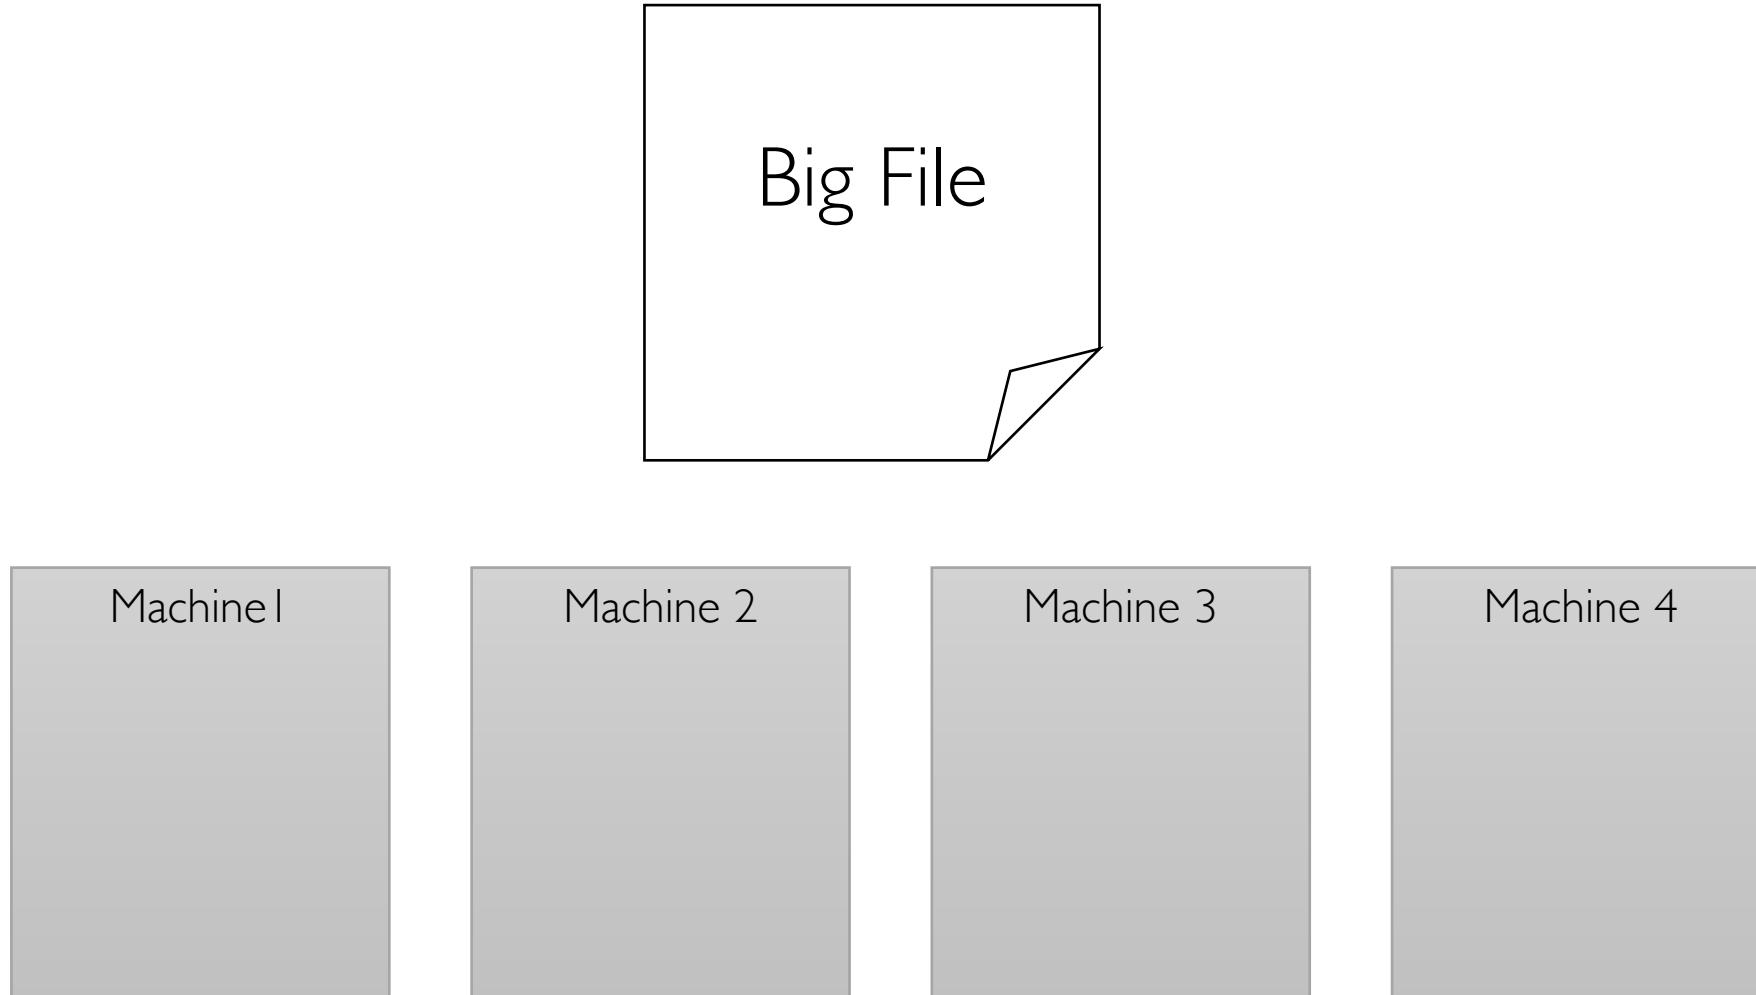
Storing very large files

Fault Tolerant Distributed File Systems

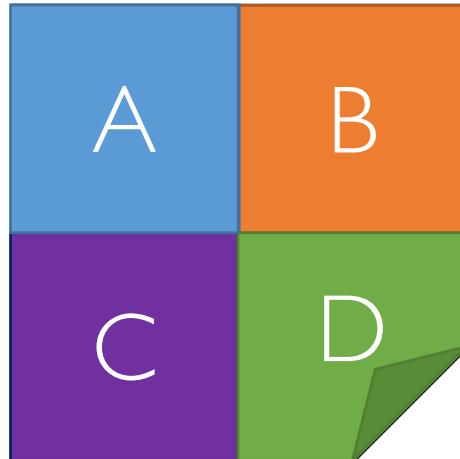


How do we **store** and **access** very
large files across **cheap**
commodity devices ?

Fault Tolerant Distributed File Systems



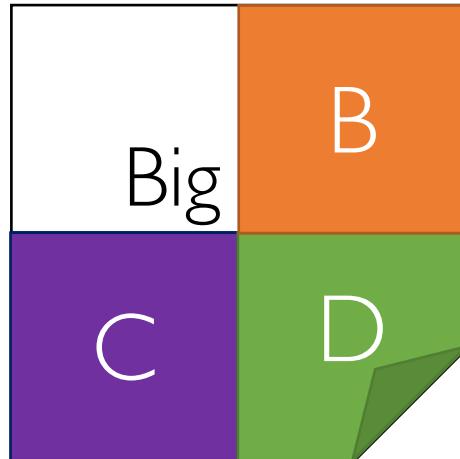
Fault Tolerant Distributed File Systems



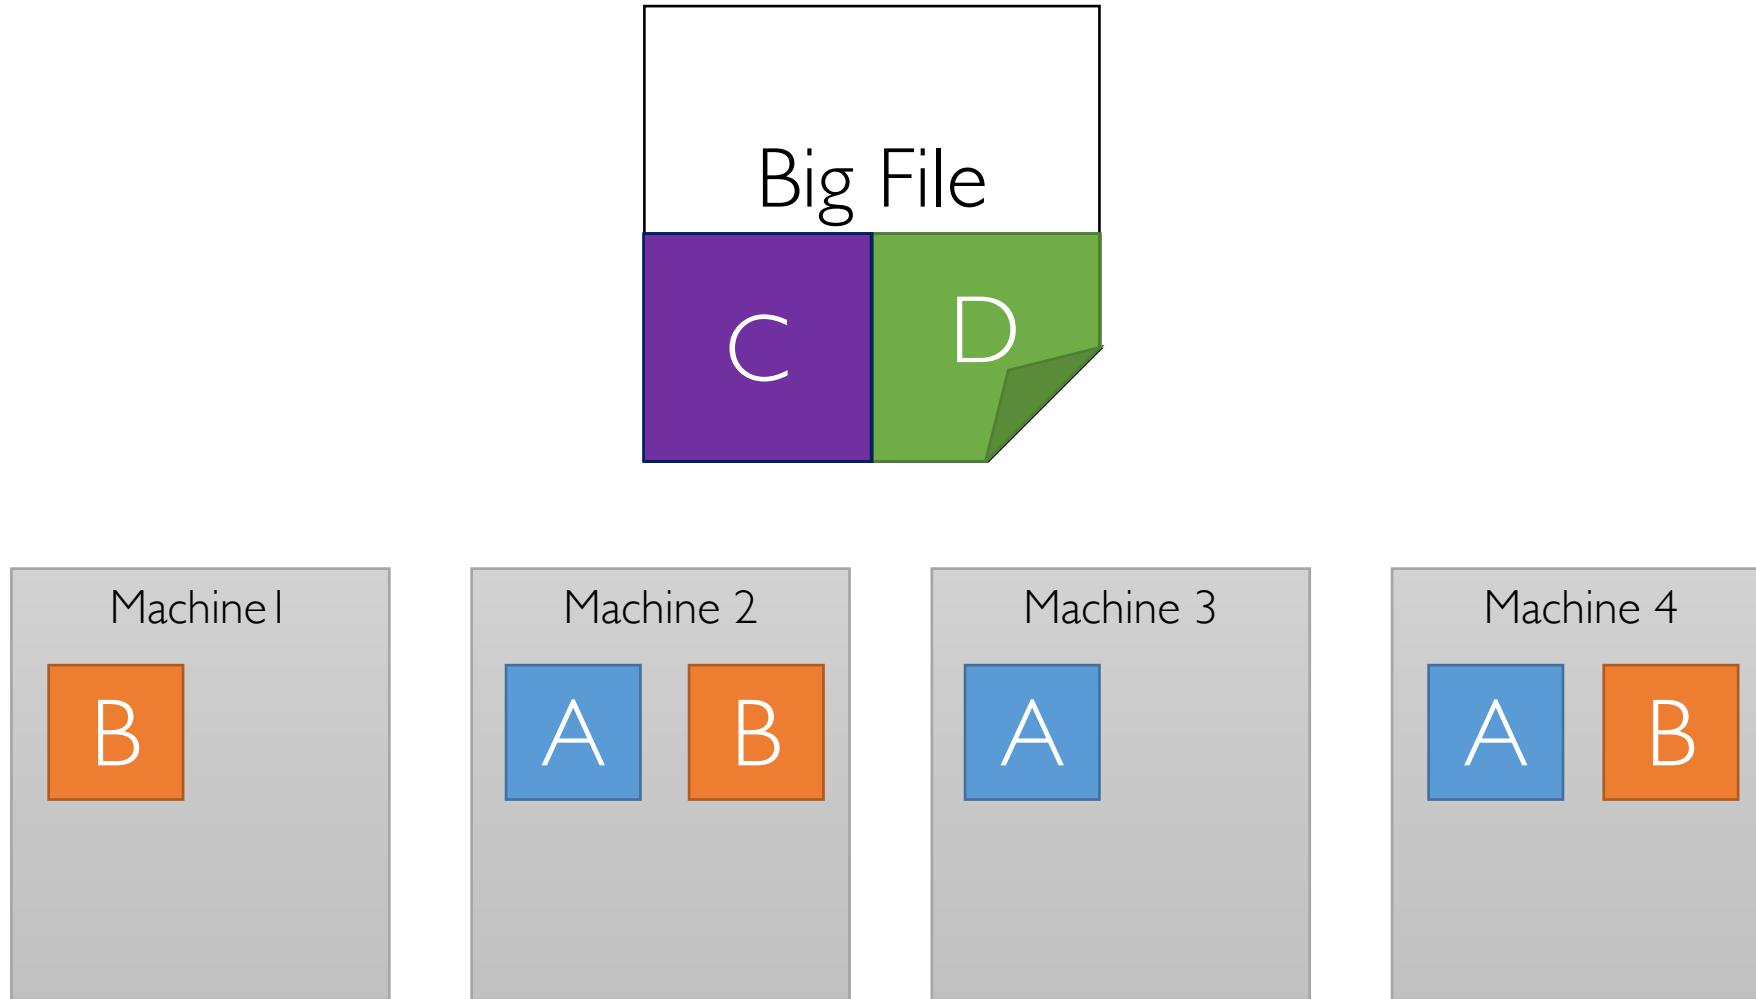
- Split the file into smaller parts.
- How?
 - Ideally at record boundaries
 - What if records are big?



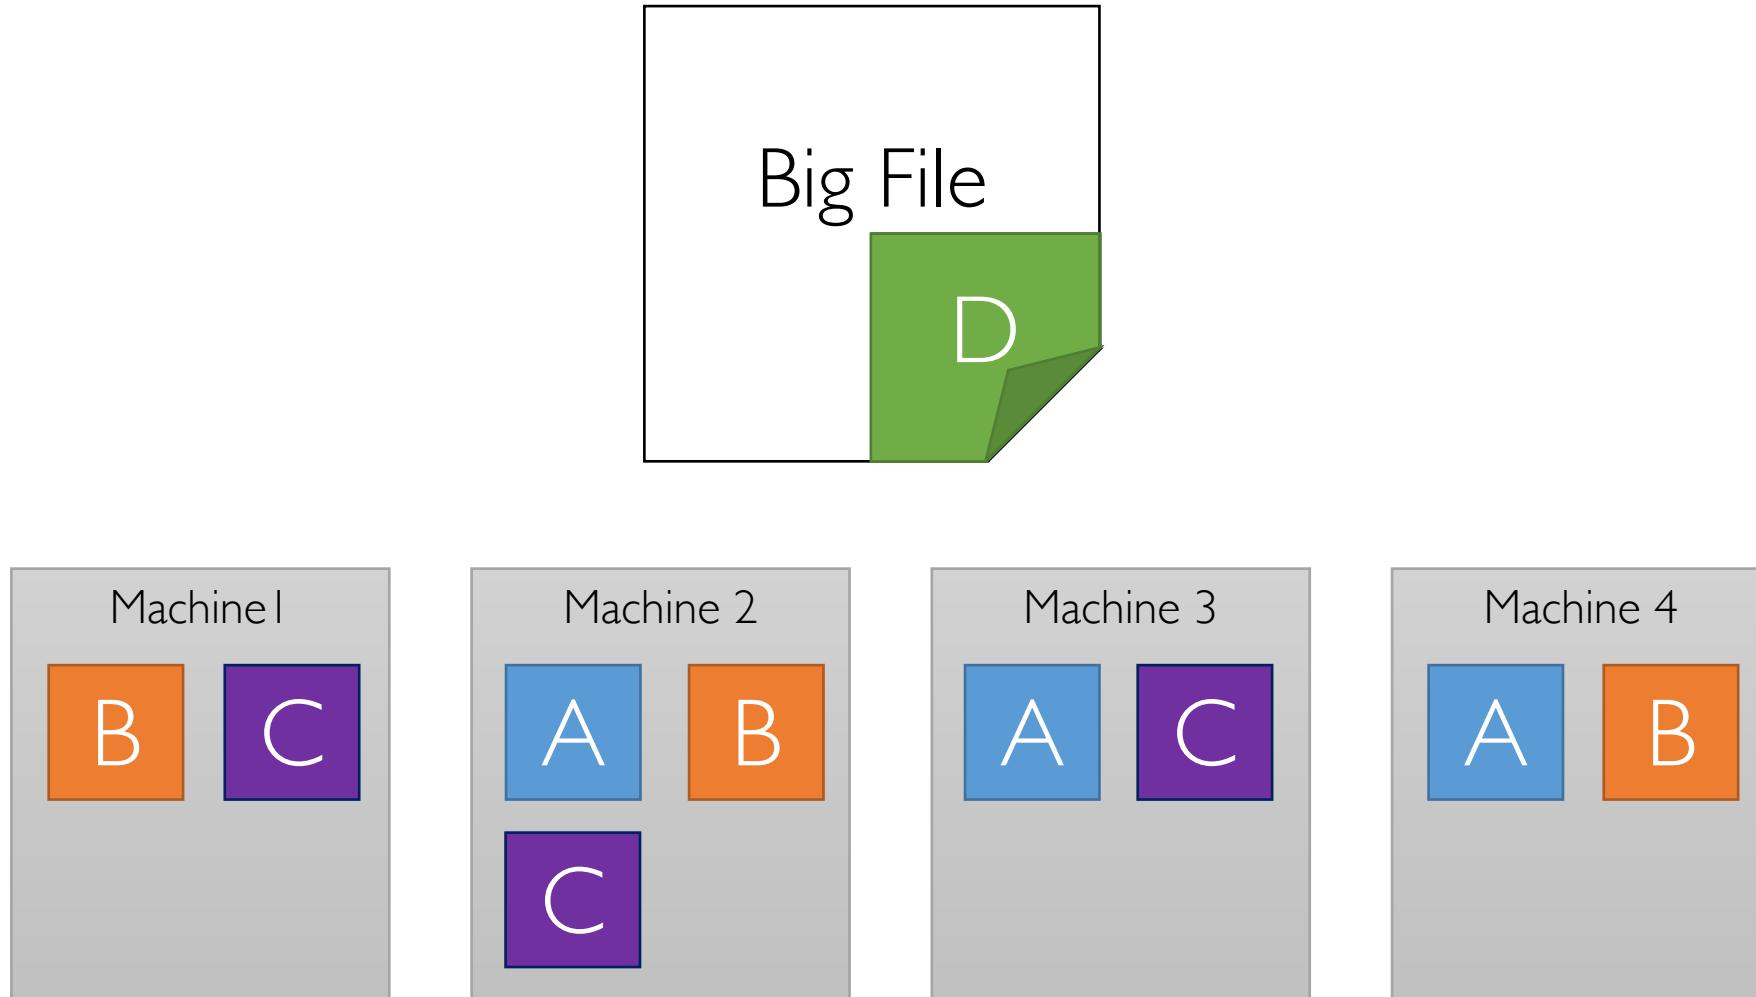
Fault Tolerant Distributed File Systems



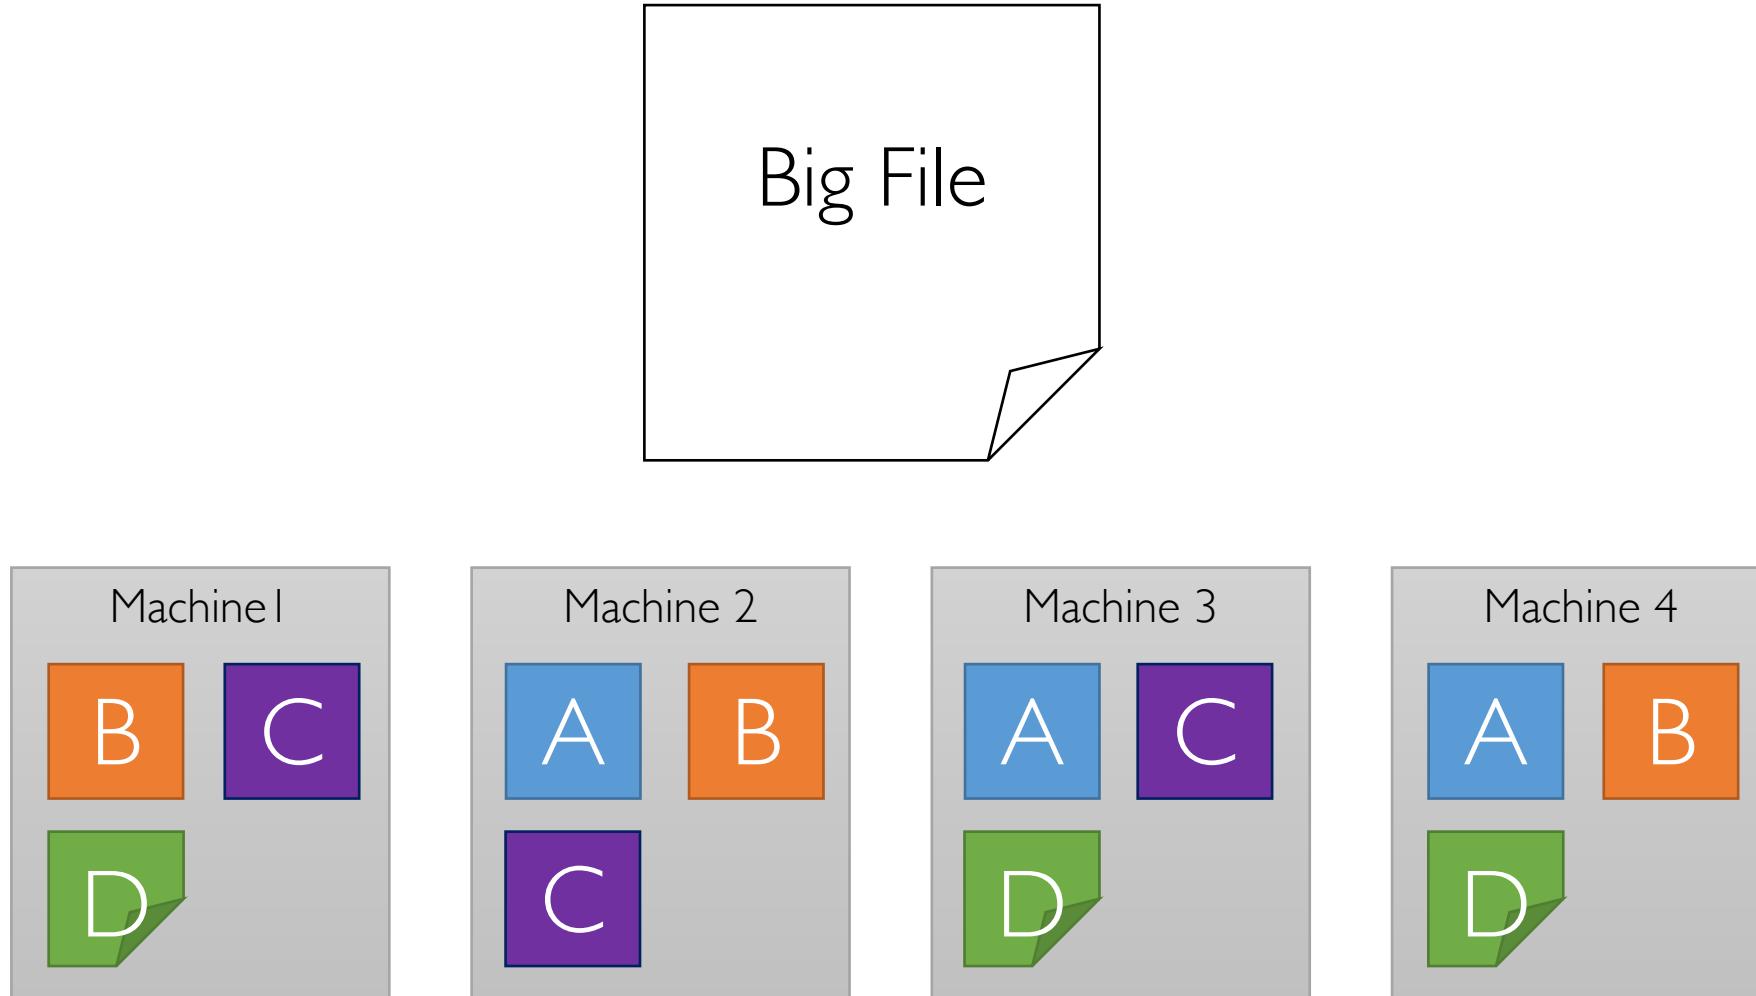
Fault Tolerant Distributed File Systems



Fault Tolerant Distributed File Systems



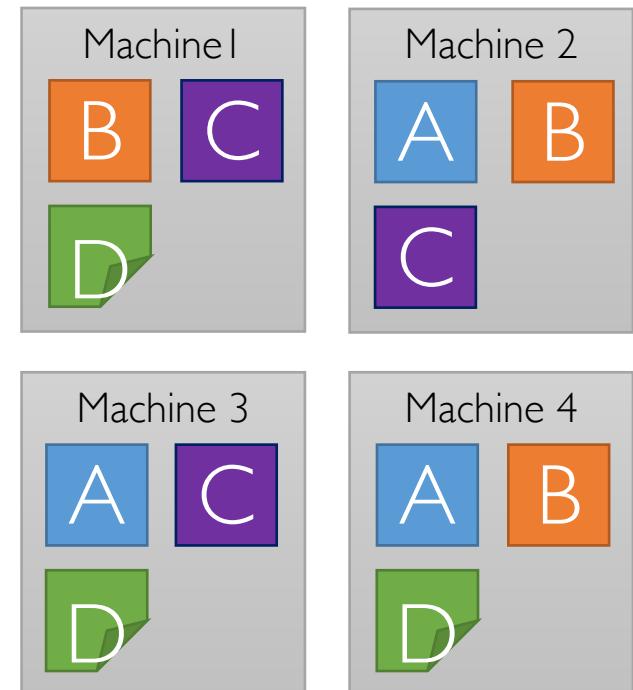
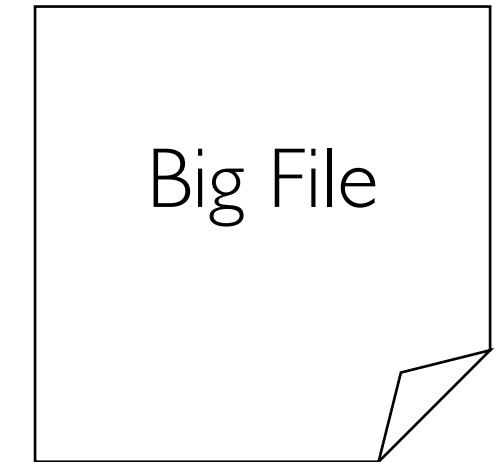
Fault Tolerant Distributed File Systems



Fault Tolerant Distributed File Systems

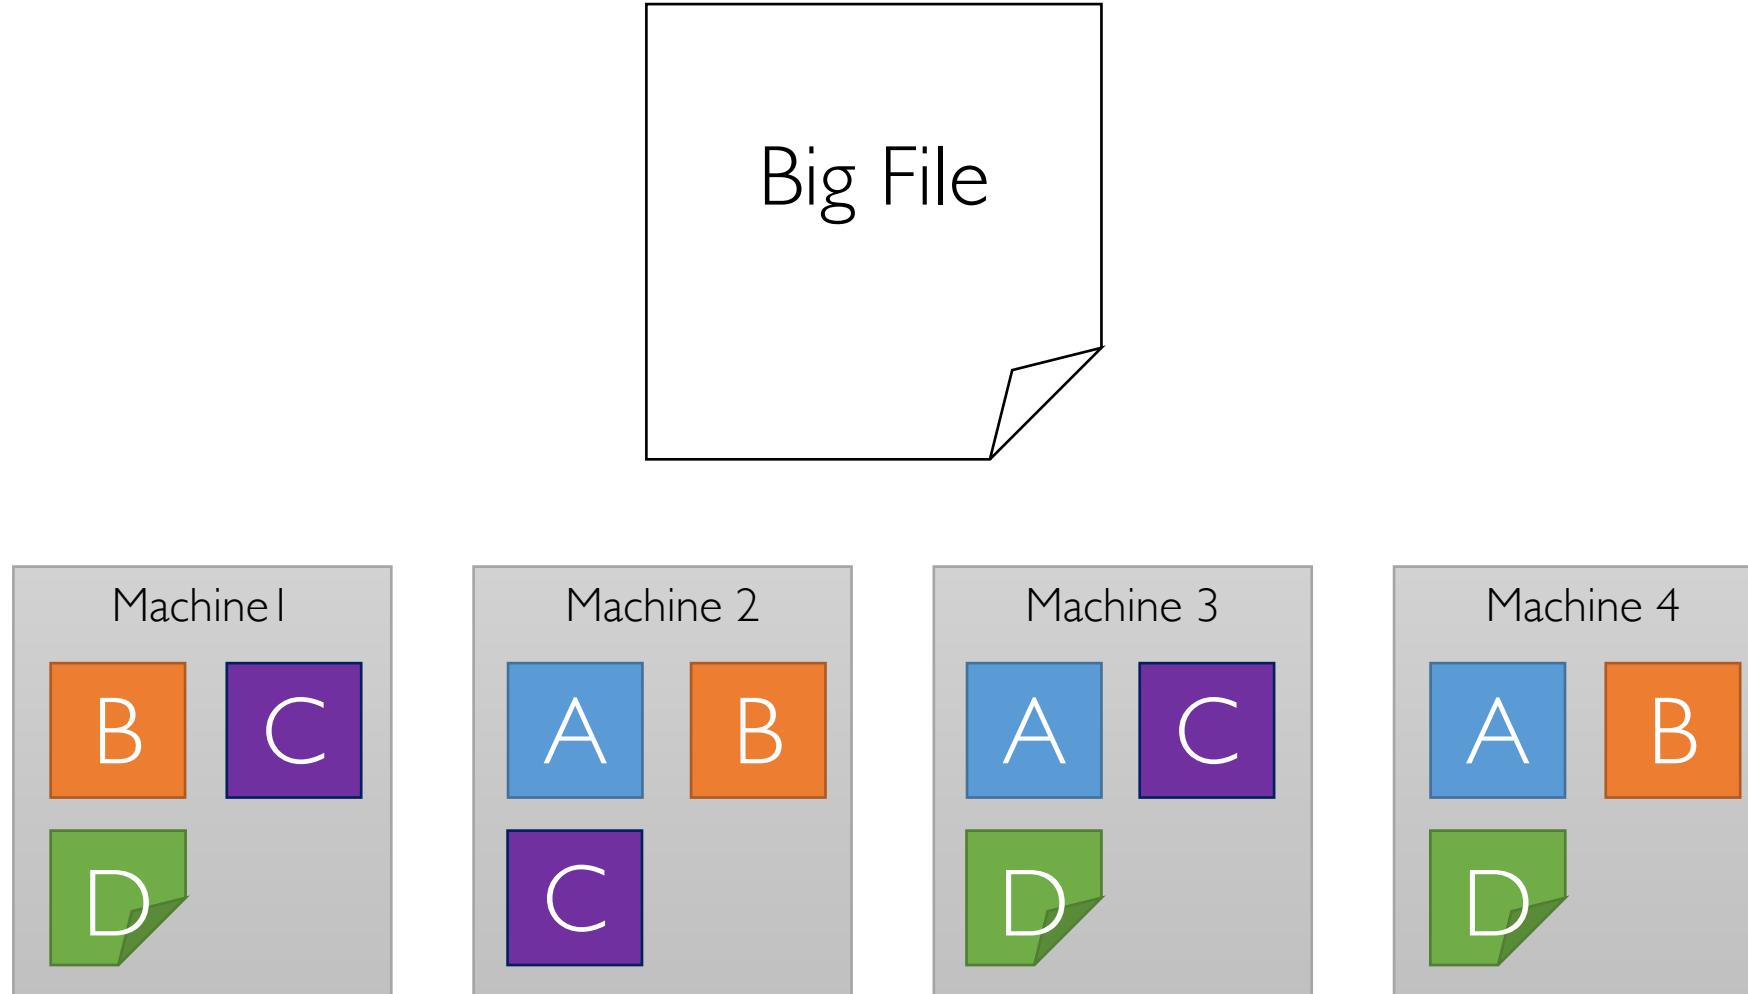
- Split large files over multiple machines
 - Easily support massive files spanning machines
- Read parts of file in parallel
 - Fast reads of large files
- Often built using cheap commodity storage devices

Cheap commodity storage devices will fail!



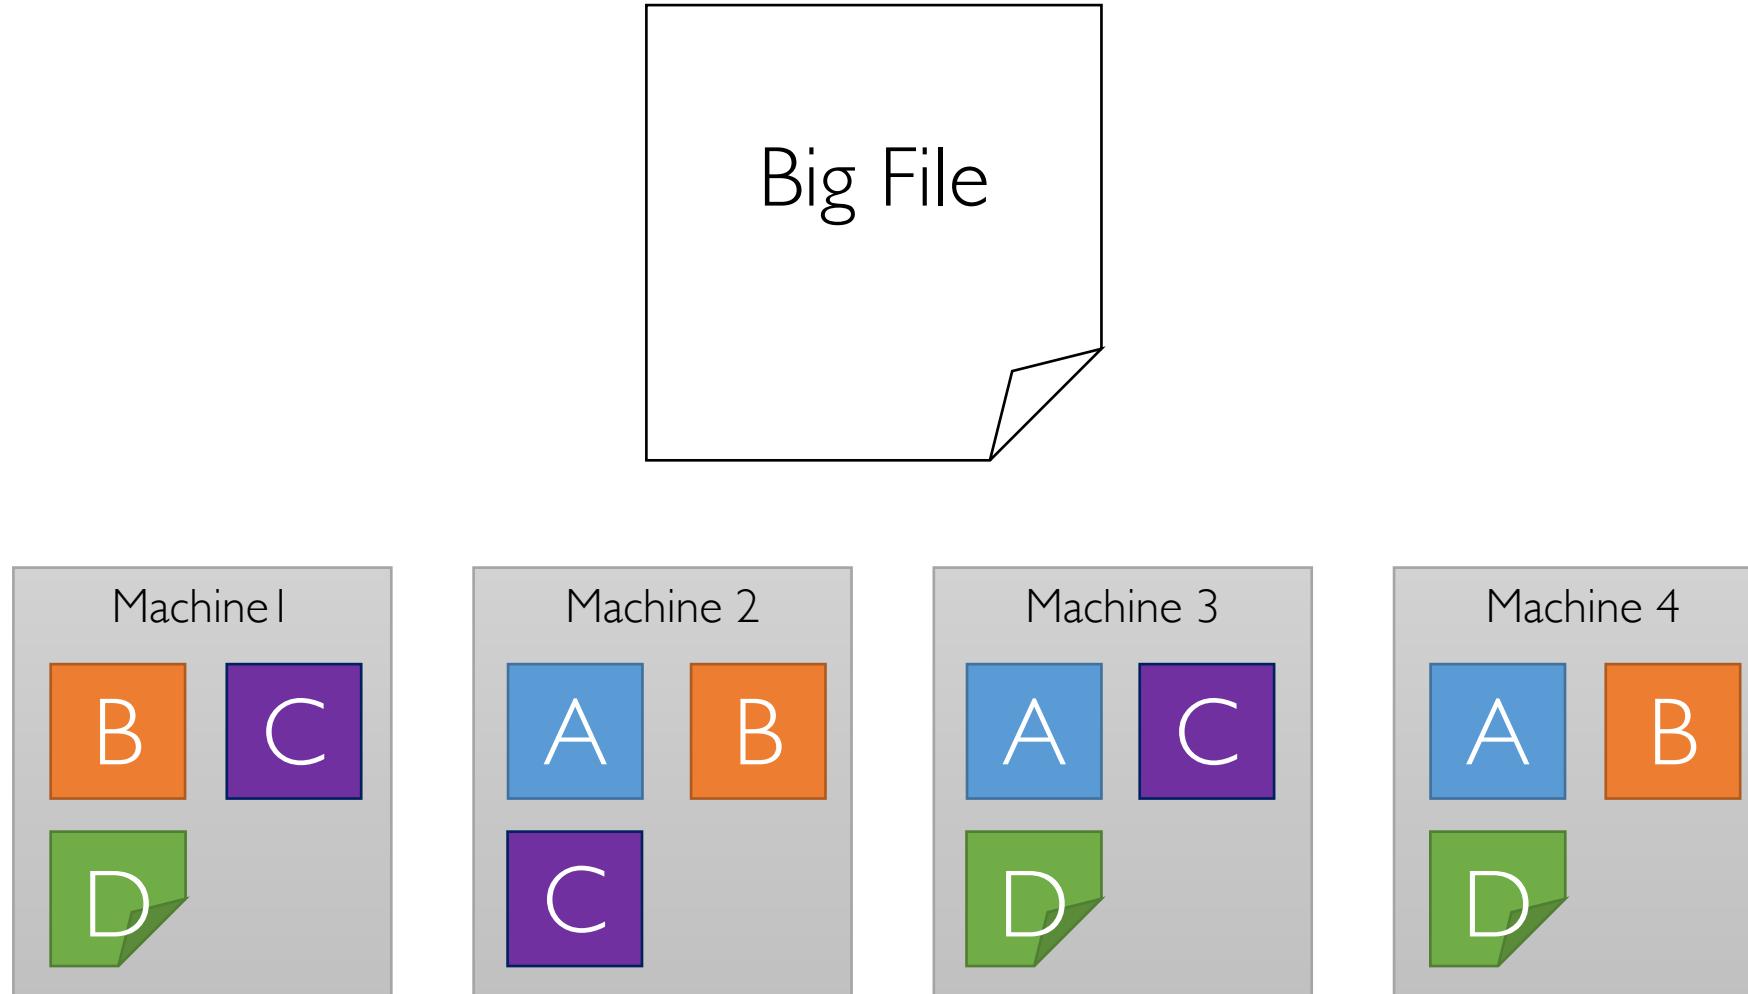
Fault Tolerant Distributed File Systems

Failure Event



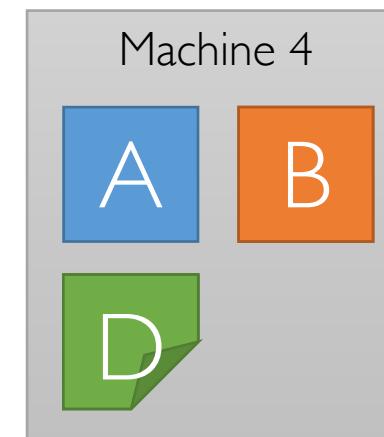
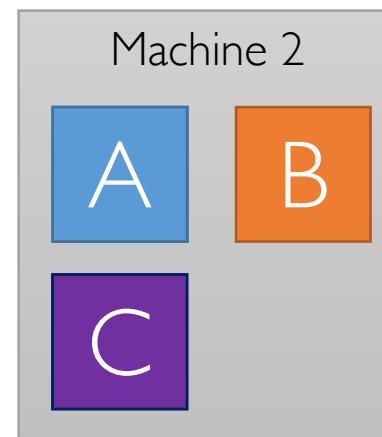
Fault Tolerant Distributed File Systems

Failure Event



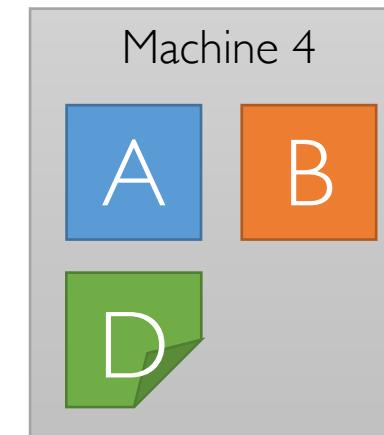
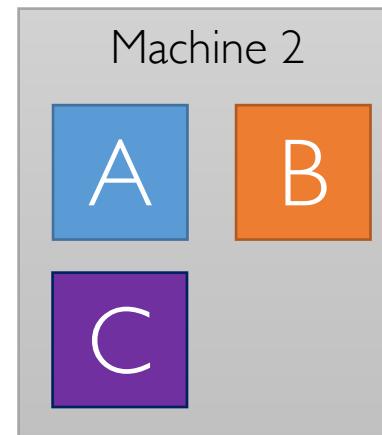
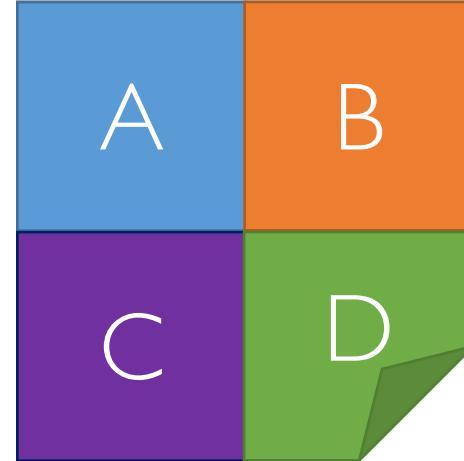
Fault Tolerant Distributed File Systems

Failure Event



Fault Tolerant Distributed File Systems

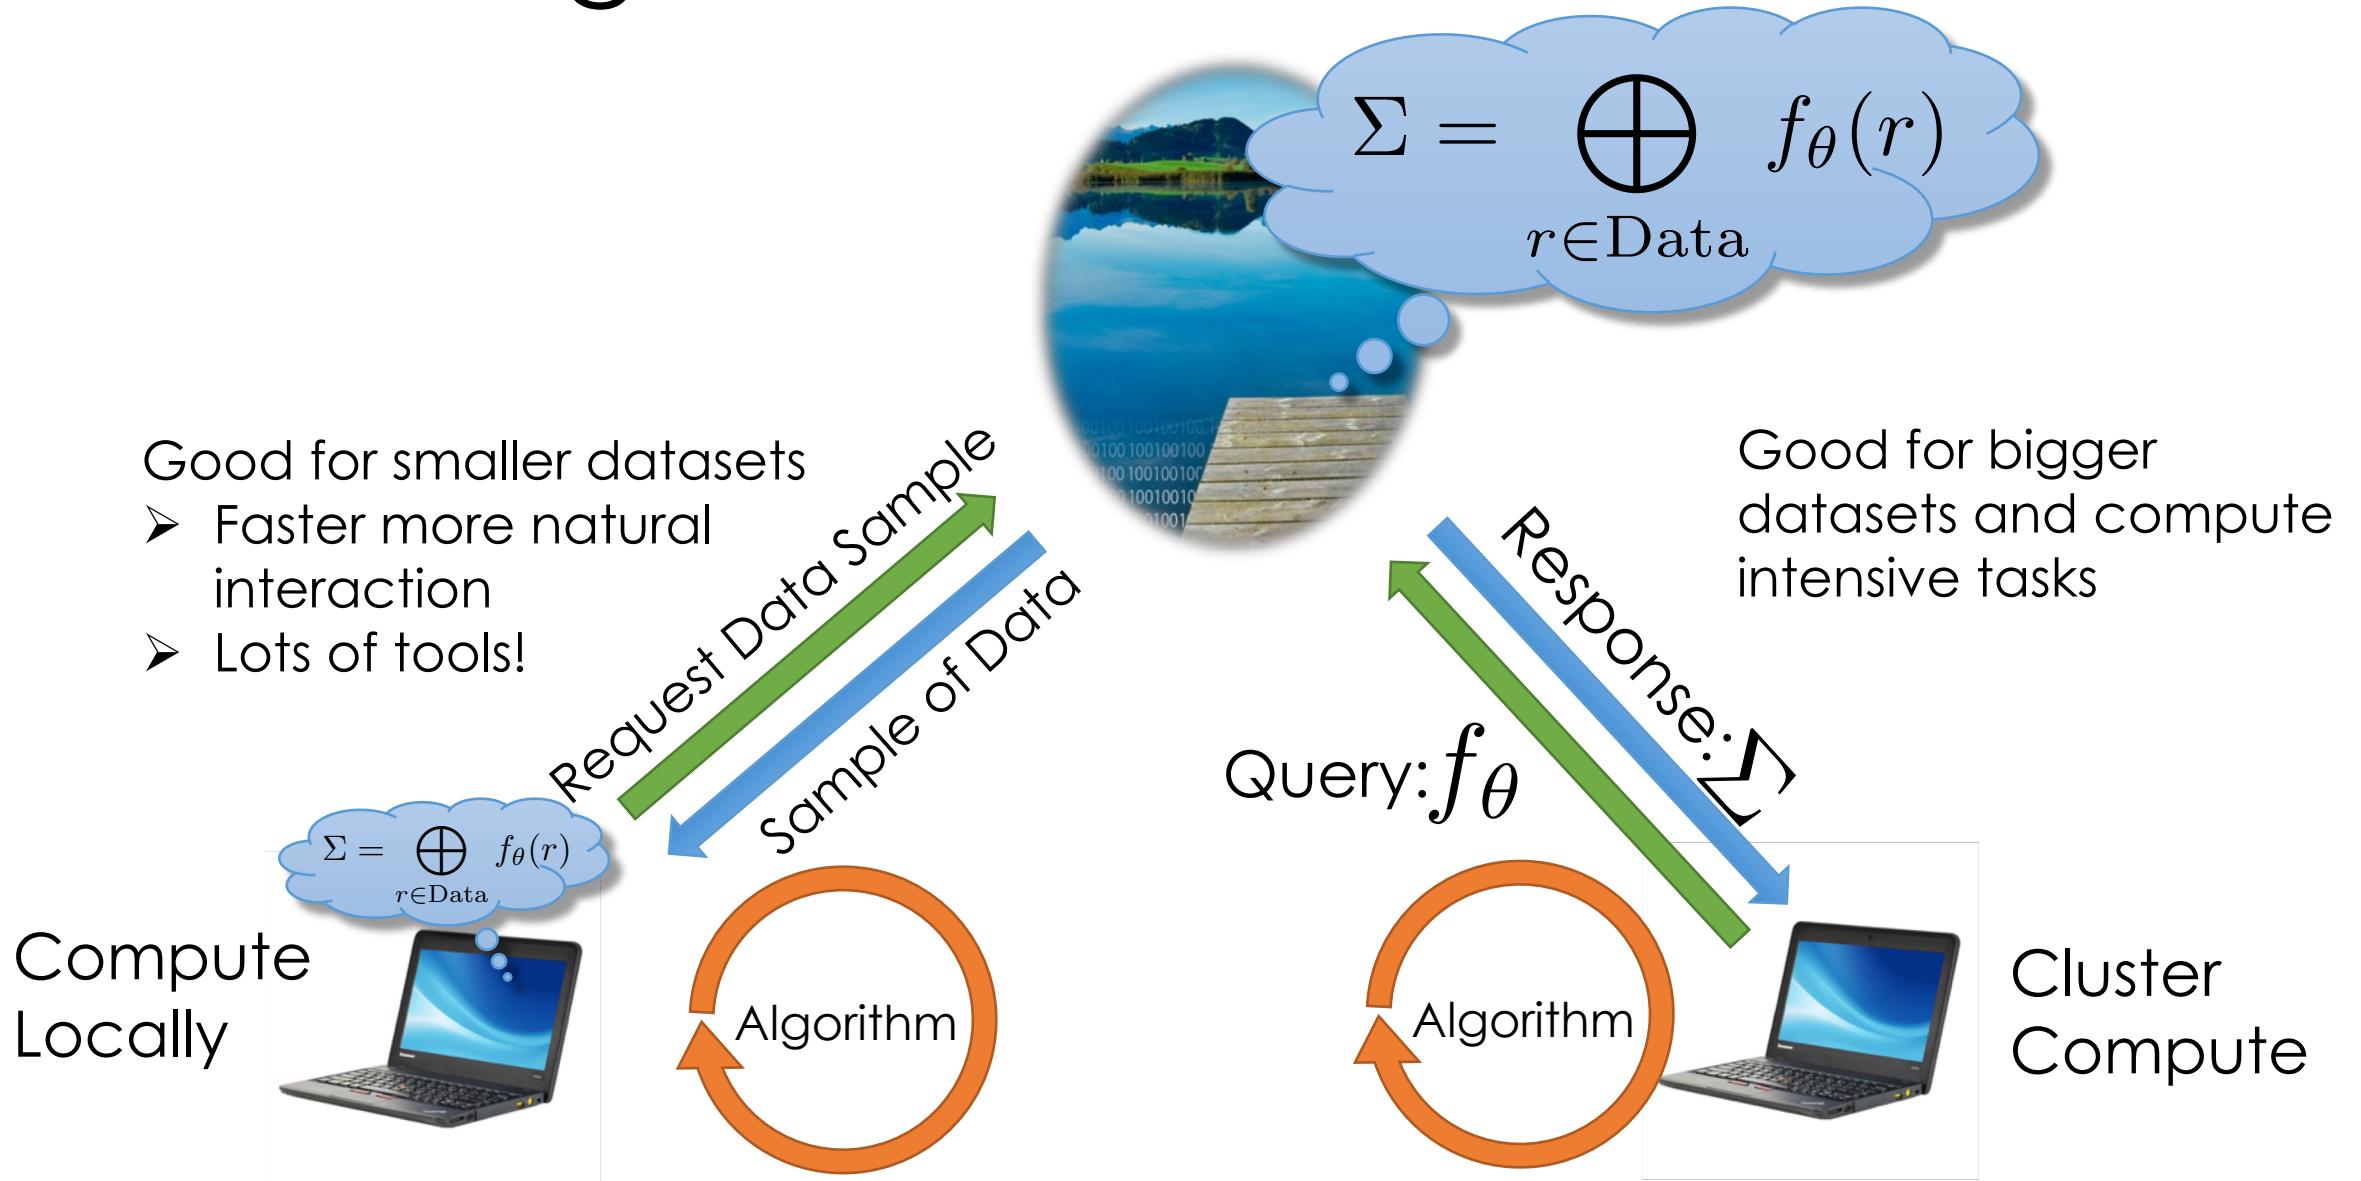
Failure Event

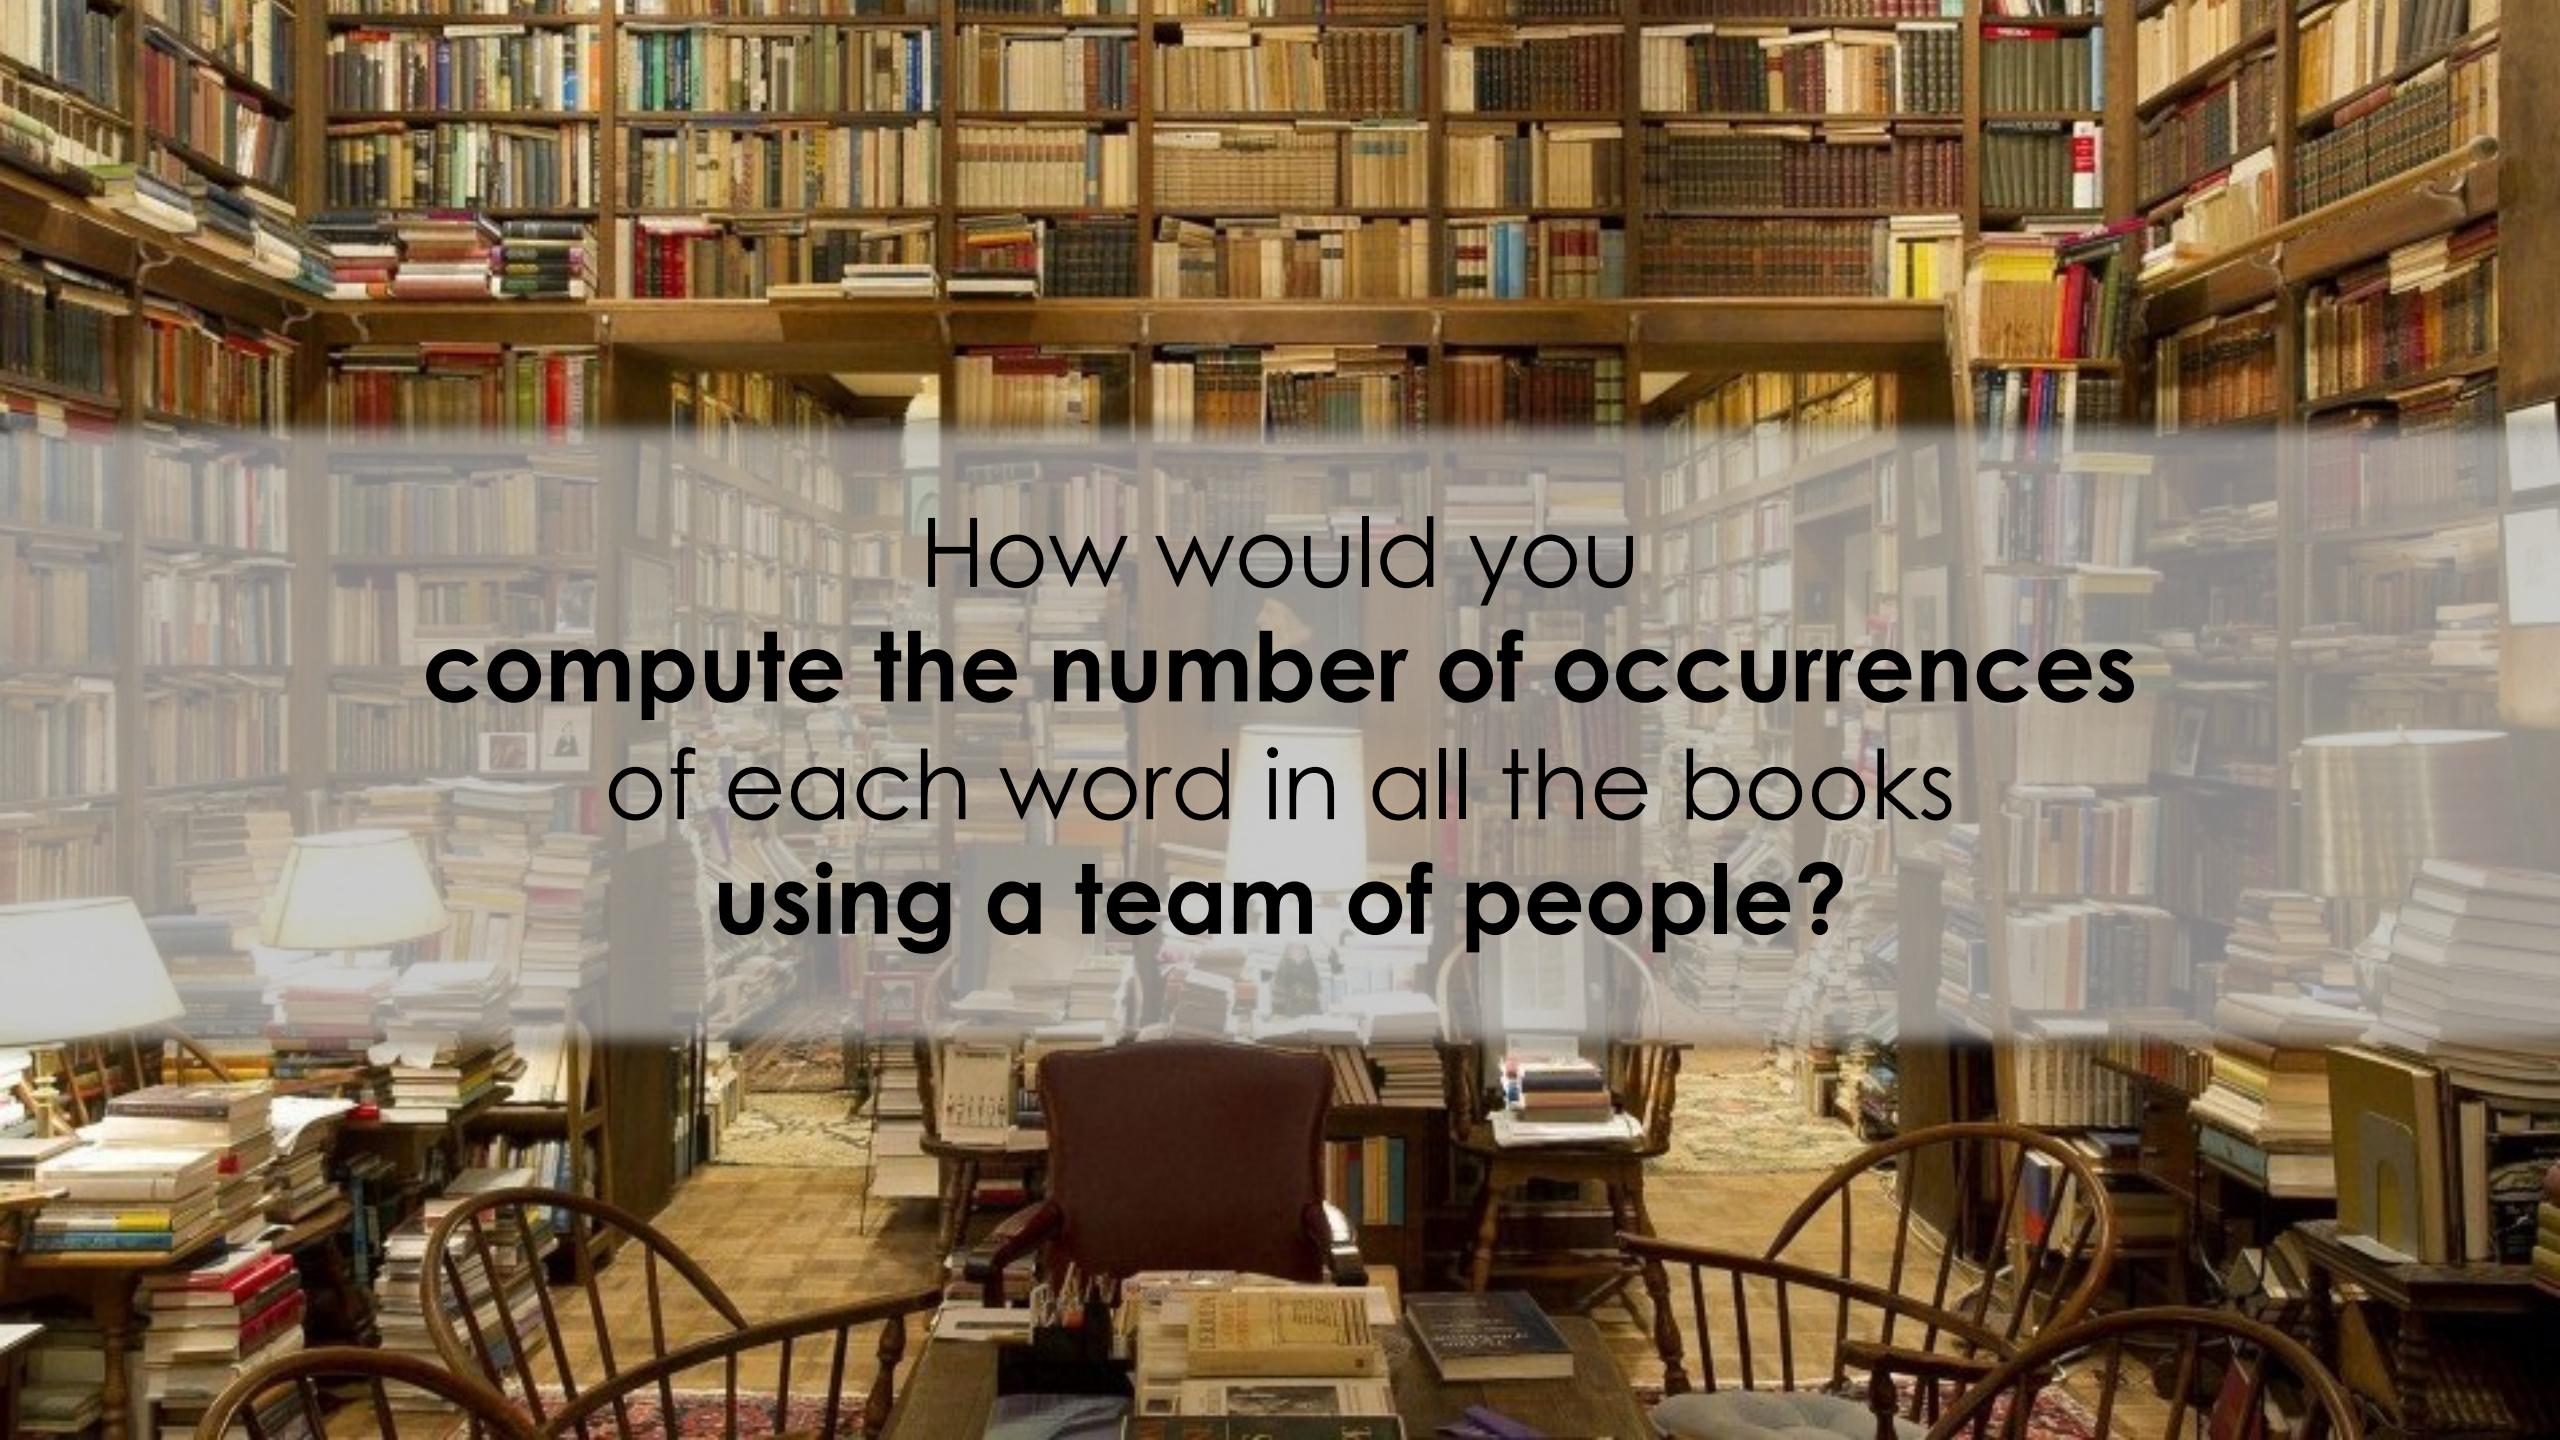


Map-Reduce Distributed Aggregation

Computing across very large files

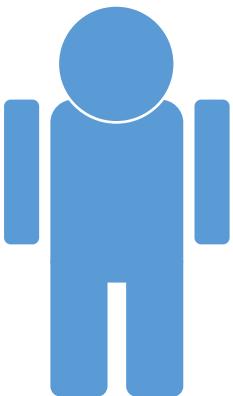
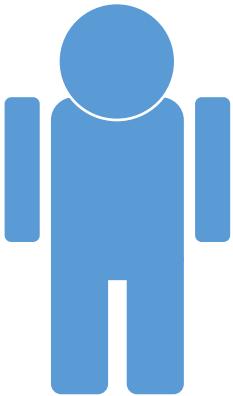
Interacting With the Data



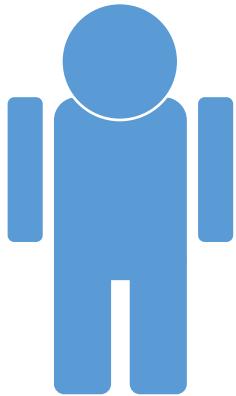


How would you
compute the number of occurrences
of each word in all the books
using a team of people?

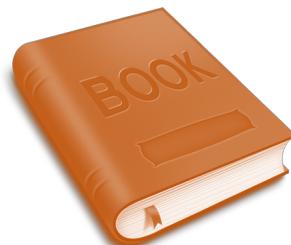
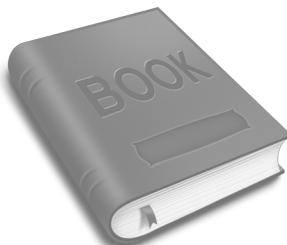
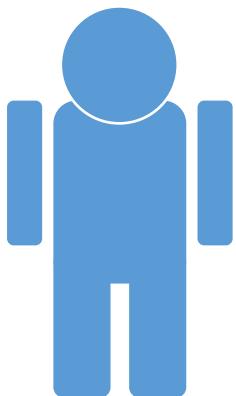
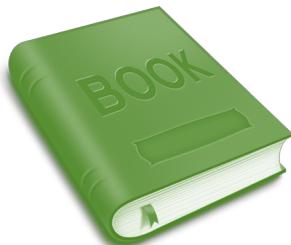
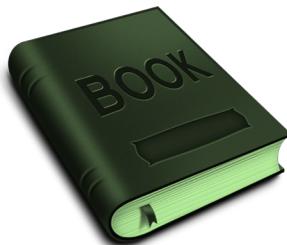
Simple Solution



Simple Solution



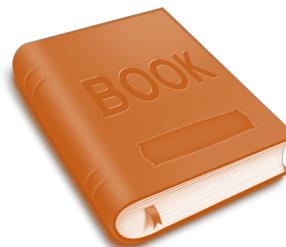
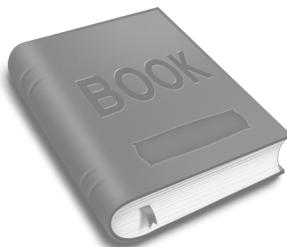
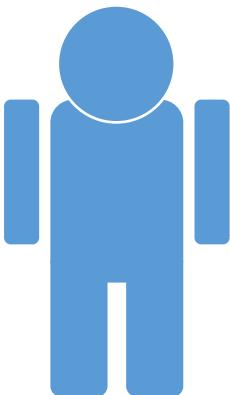
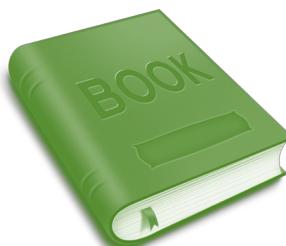
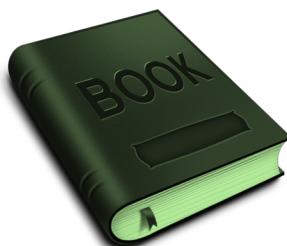
1) Divide Books Across Individuals



Simple Solution



1) Divide Books Across Individuals

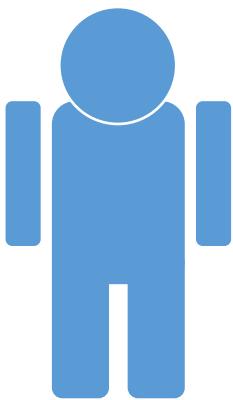


2) Compute Counts Locally

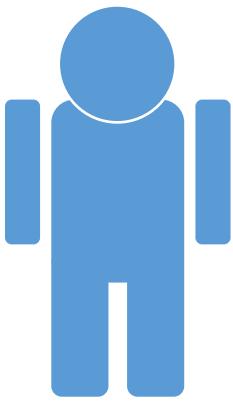
Word	Count
Apple	2
Bird	7
...	

Word	Count
Apple	0
Bird	1
...	

Simple Solution



1) Divide Books Across Individuals



2) Compute Counts Locally

Word	Count
Apple	2
Bird	7
...	

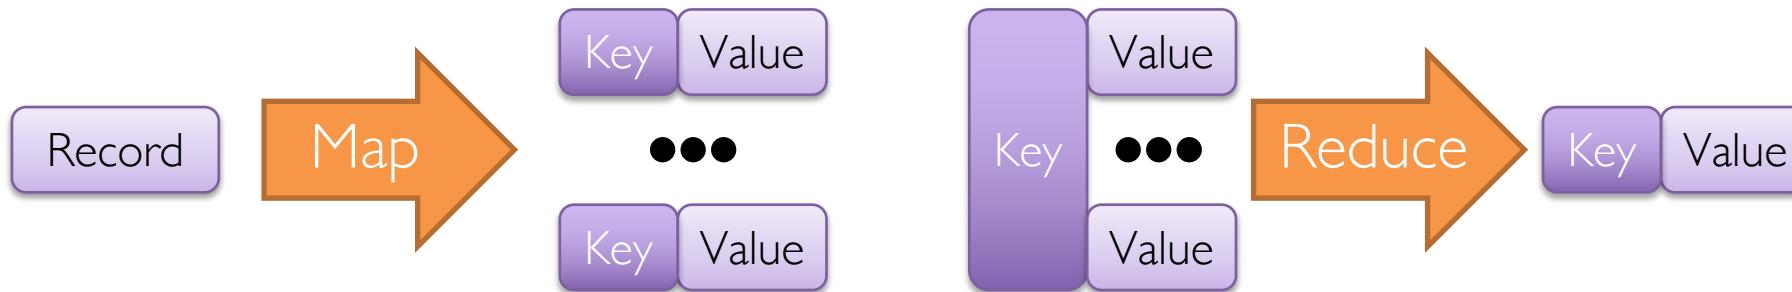
3) Aggregate Tables



Word	Count
Apple	2
Bird	8
...	

Word	Count
Apple	0
Bird	1
...	

The Map Reduce Abstraction



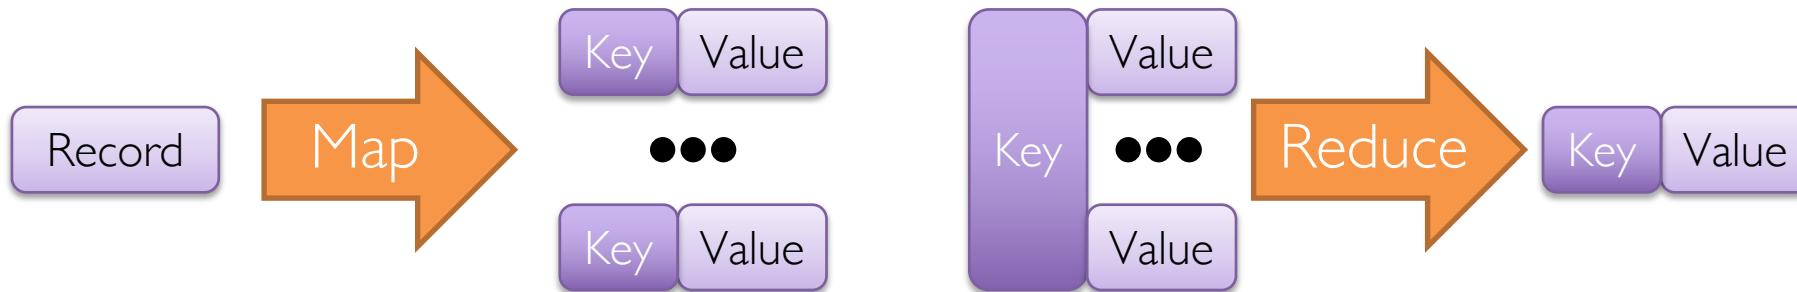
Example: Word-Count

```
Map(book):  
for (word in set(book)):  
    emit (word, book.count(word))
```

Key Value

```
Reduce(word, counts) {  
    sum = 0  
    for count in counts:  
        sum += count  
    emit (word, SUM(counts))  
}
```

The Map Reduce Abstraction (Simpler)



Example: Word-Count

Map(book):
for (word in book):
 emit (word, 1)

Key Value

```
Reduce(word, counts) {  
    sum = 0  
    for count in counts:  
        sum += count  
    emit (word, SUM(counts))  
}
```

The Map Reduce Abstraction (General)



Example: Word-Count

```
Map(record, f):  
    for (key in record):  
        emit (key, f(key))
```

Key Value

```
Reduce(key, values, f) {  
    agg = f(values[0], values[1])  
    for value in values[2:]:  
        agg = f(agg, value)  
    emit (word, agg)  
}
```

Map: Deterministic

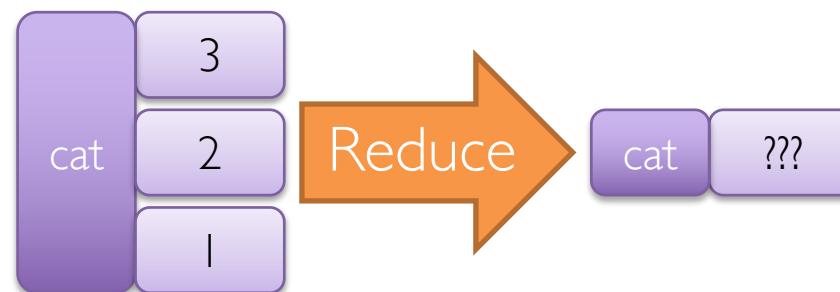
Reduce: Commutative and Associative

Key properties of Map And Reduce

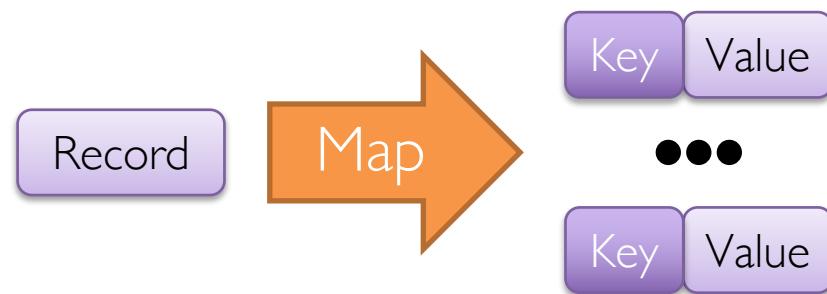
- **Deterministic Map:** allows for re-execution on failure
 - If some computation is lost we can always re-compute
 - Issues with samples?
- **Commutative Reduce:** allows for re-order of operations
 - $\text{Reduce}(A,B) = \text{Reduce}(B,A)$
 - Example (addition): $A + B = B + A$
- **Associative Reduce:** allows for regrouping of operations
 - $\text{Reduce}(\text{Reduce}(A,B), C) = \text{Reduce}(A, \text{Reduce}(B,C))$
 - Example (max): $\text{max}(\text{max}(A,B), C) = \text{max}(A, \text{max}(B,C))$
 - Warning: Floating point operations (e.g. addition) are not guaranteed associative.

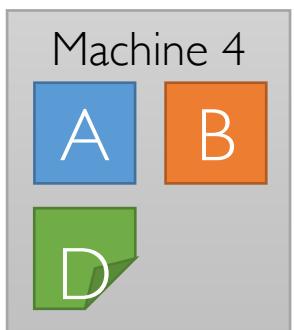
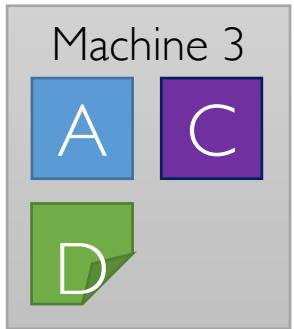
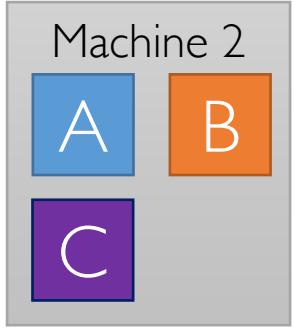
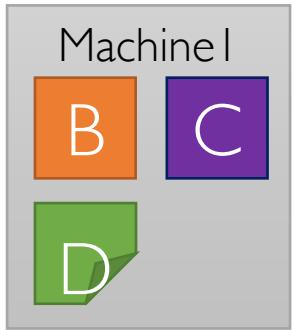
Question

- Suppose our reduction function computes $a^*b + 1$.
- Suppose we have 3 values associated with the key 'cat'. What is the result of the reduction operation?



Executing Map Reduce

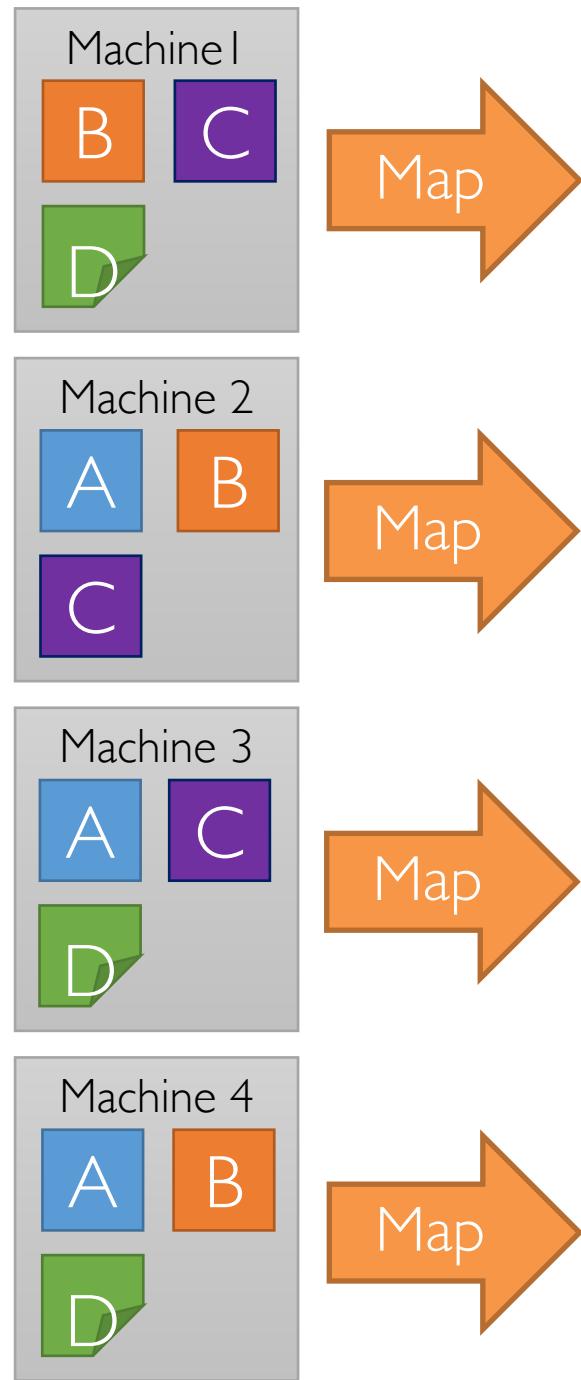




Executing Map Reduce



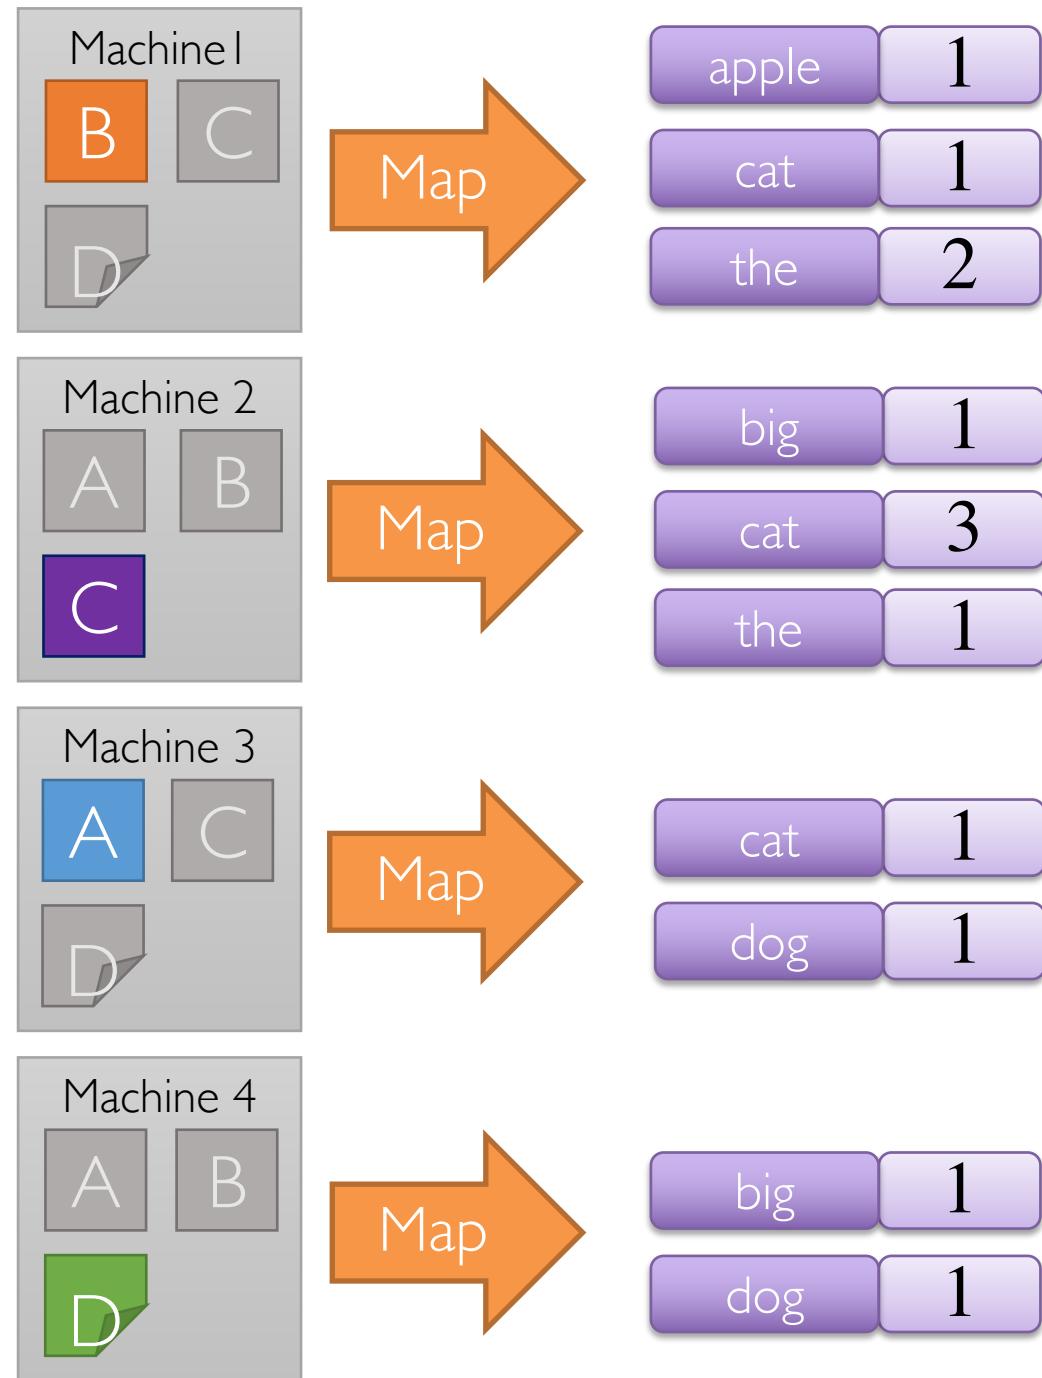
Distributing the Map Function



Executing Map Reduce

Distributing the Map Function

Executing Map Reduce

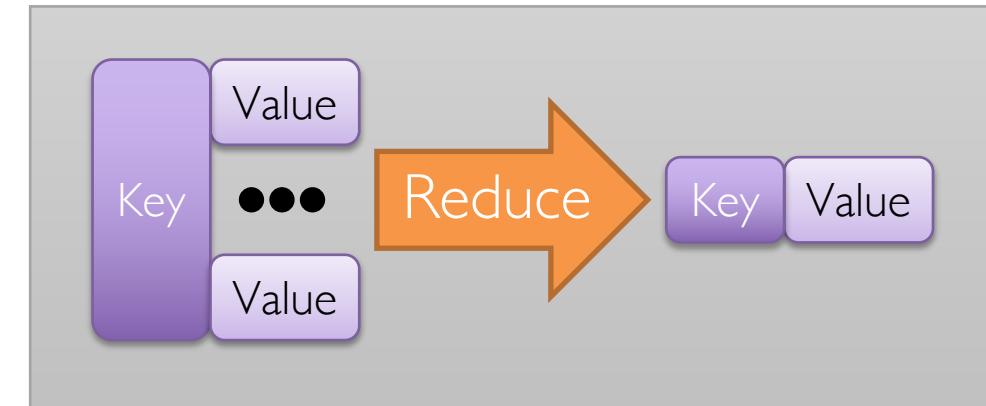
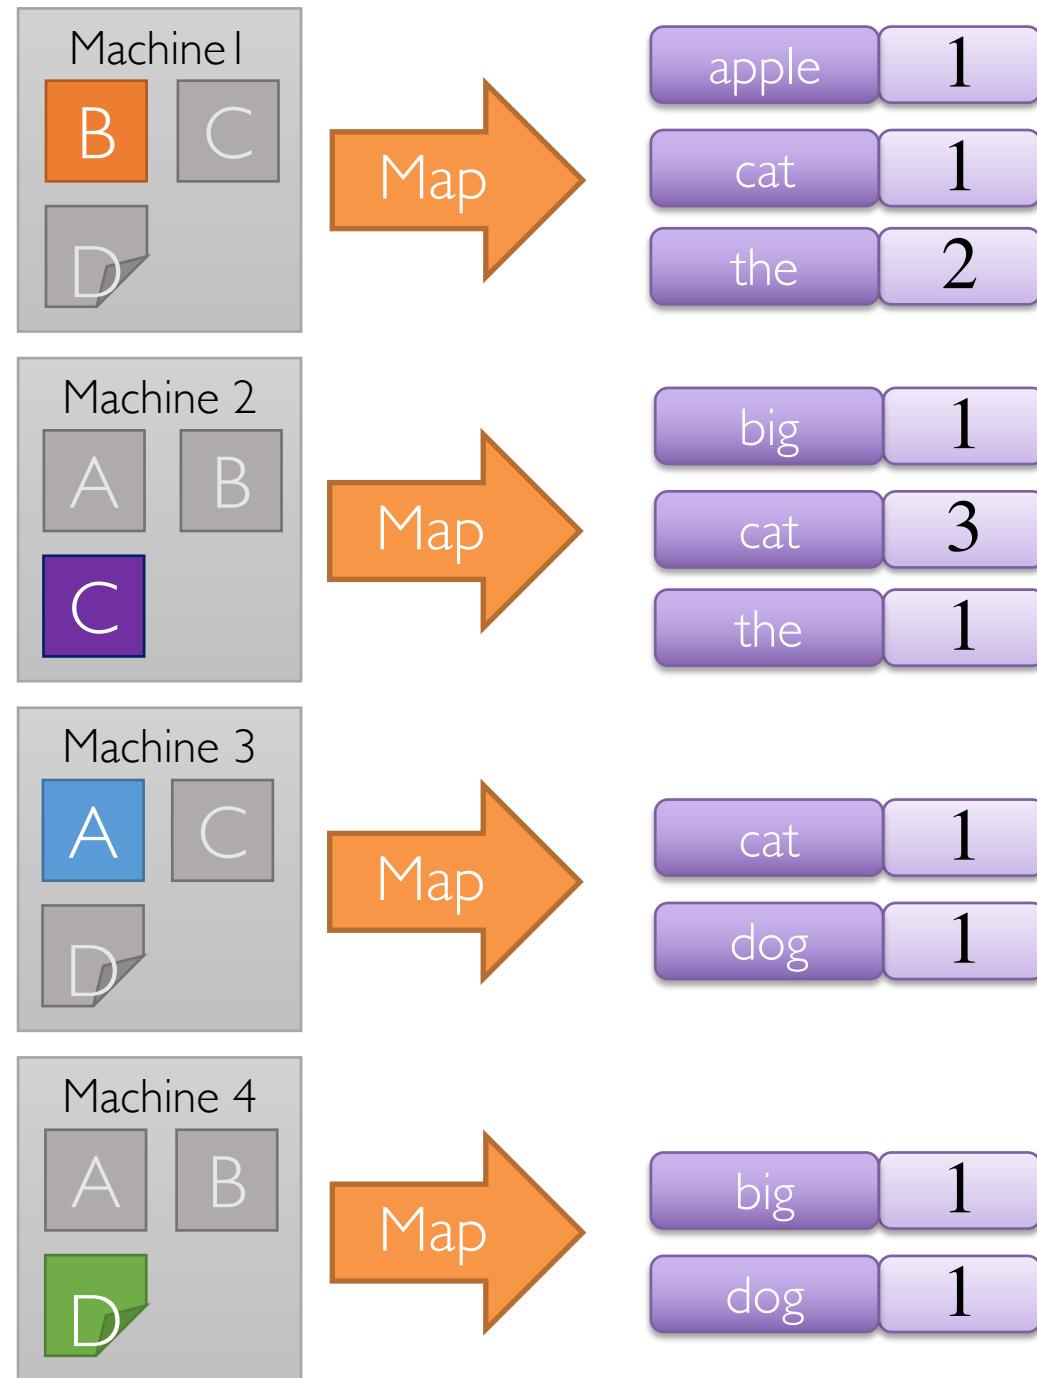


The map function applied to a local part of the big file.

Run in Parallel.

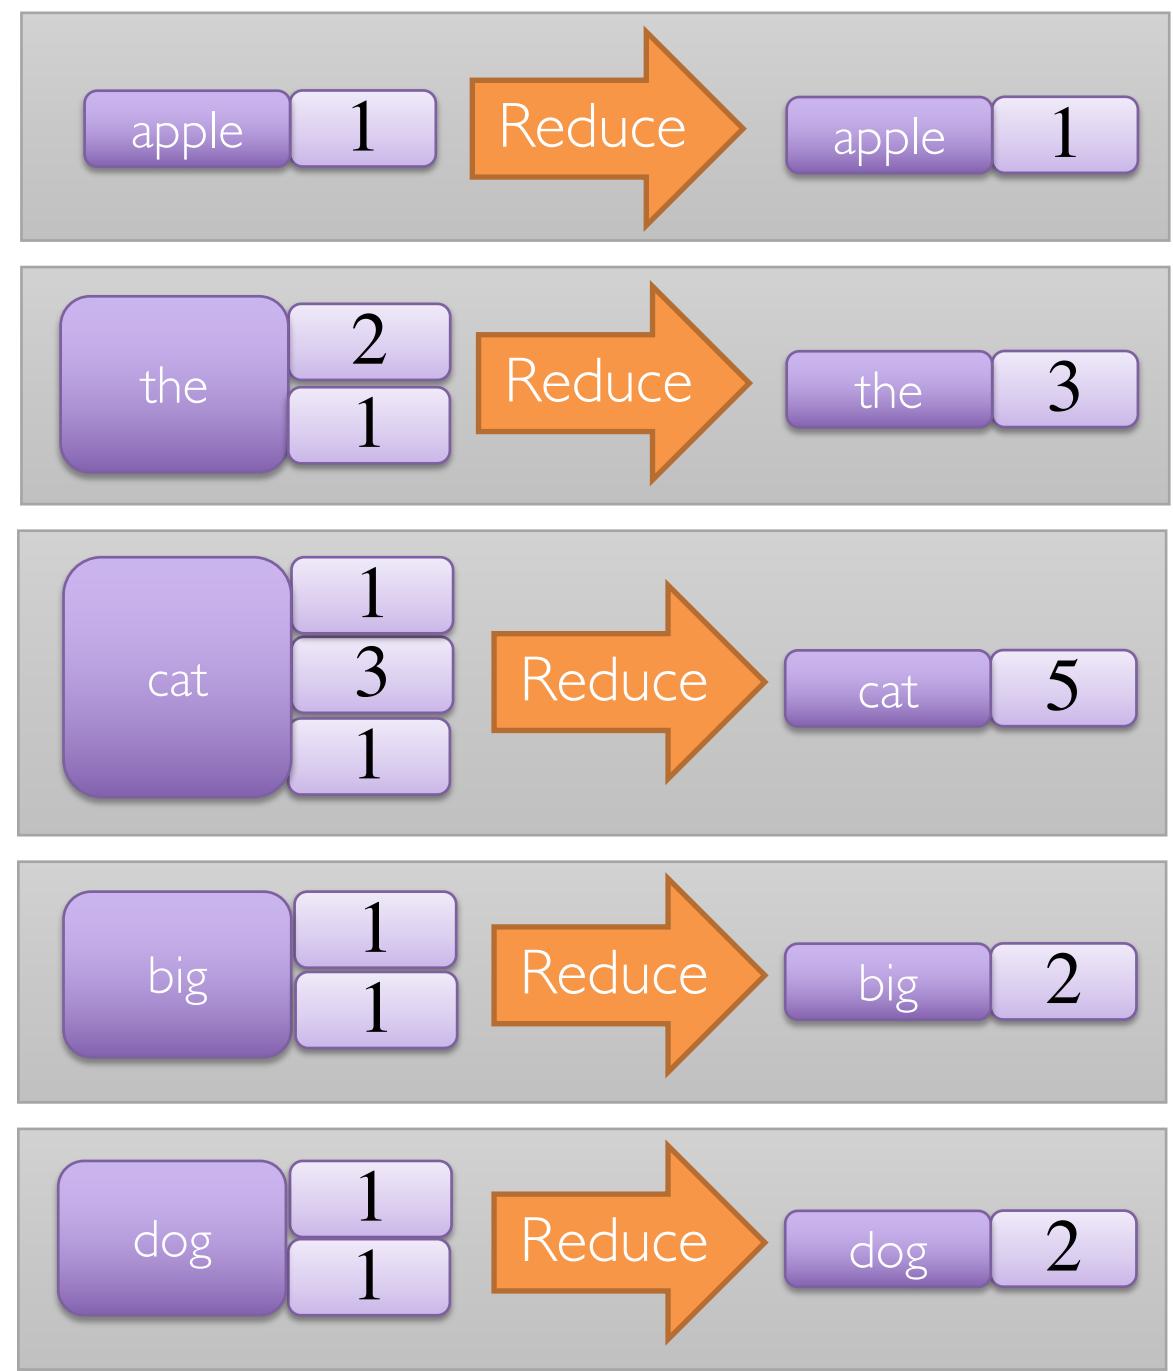
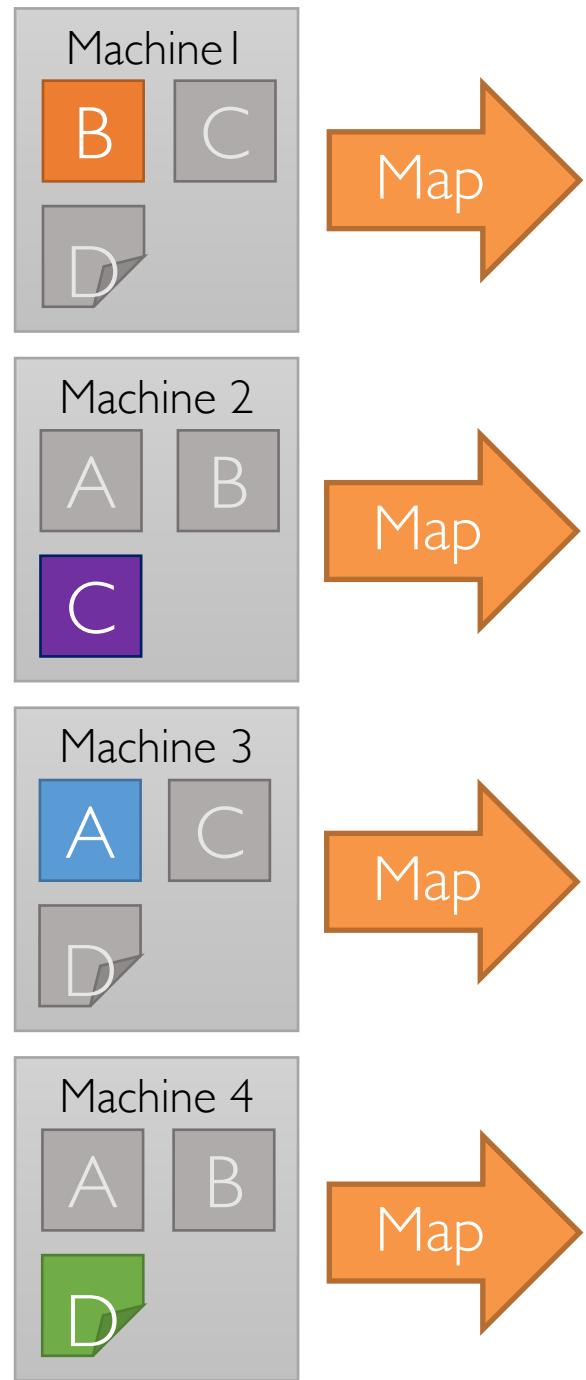
Output is cached for fast recovery on node failure

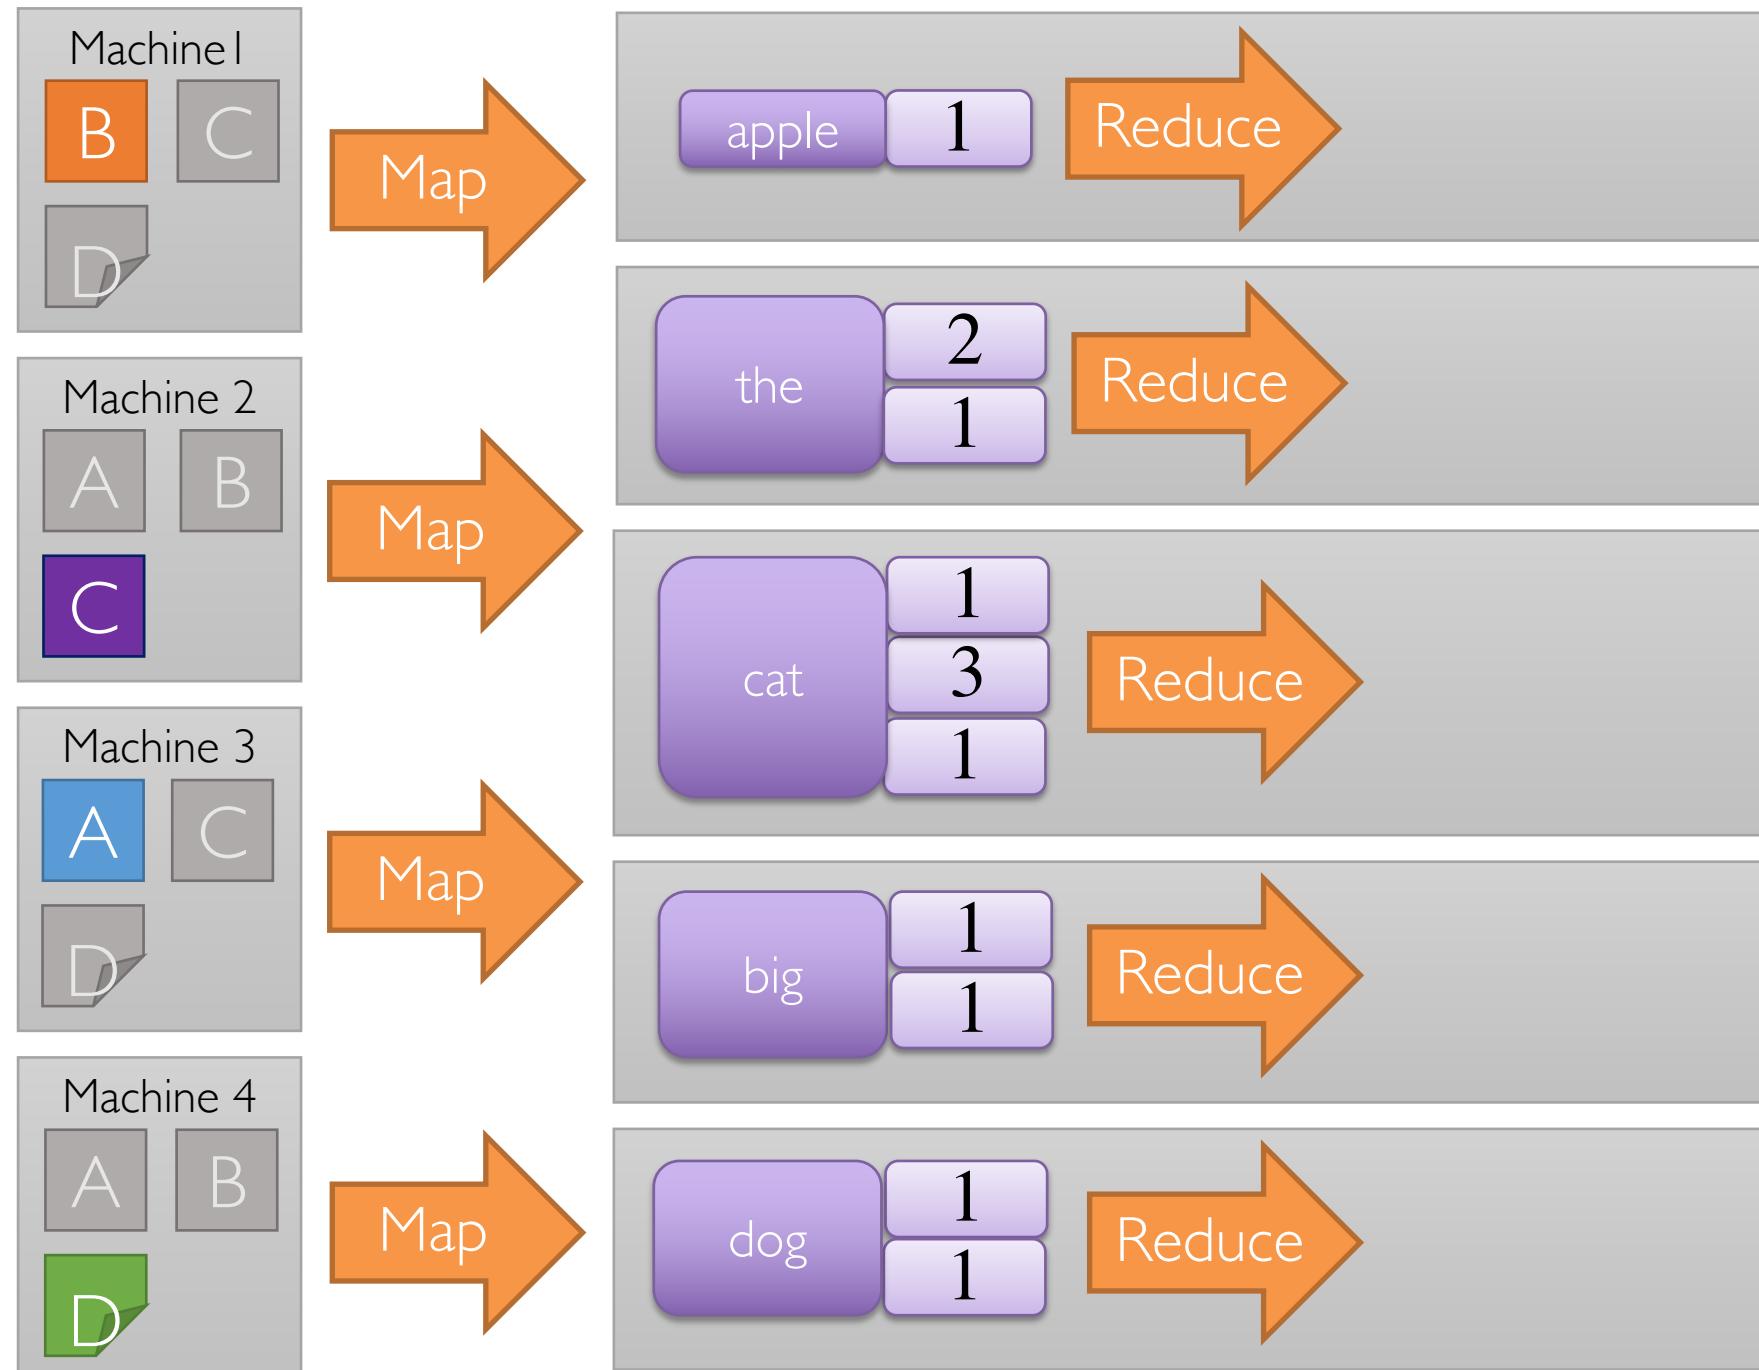
Executing Map Reduce

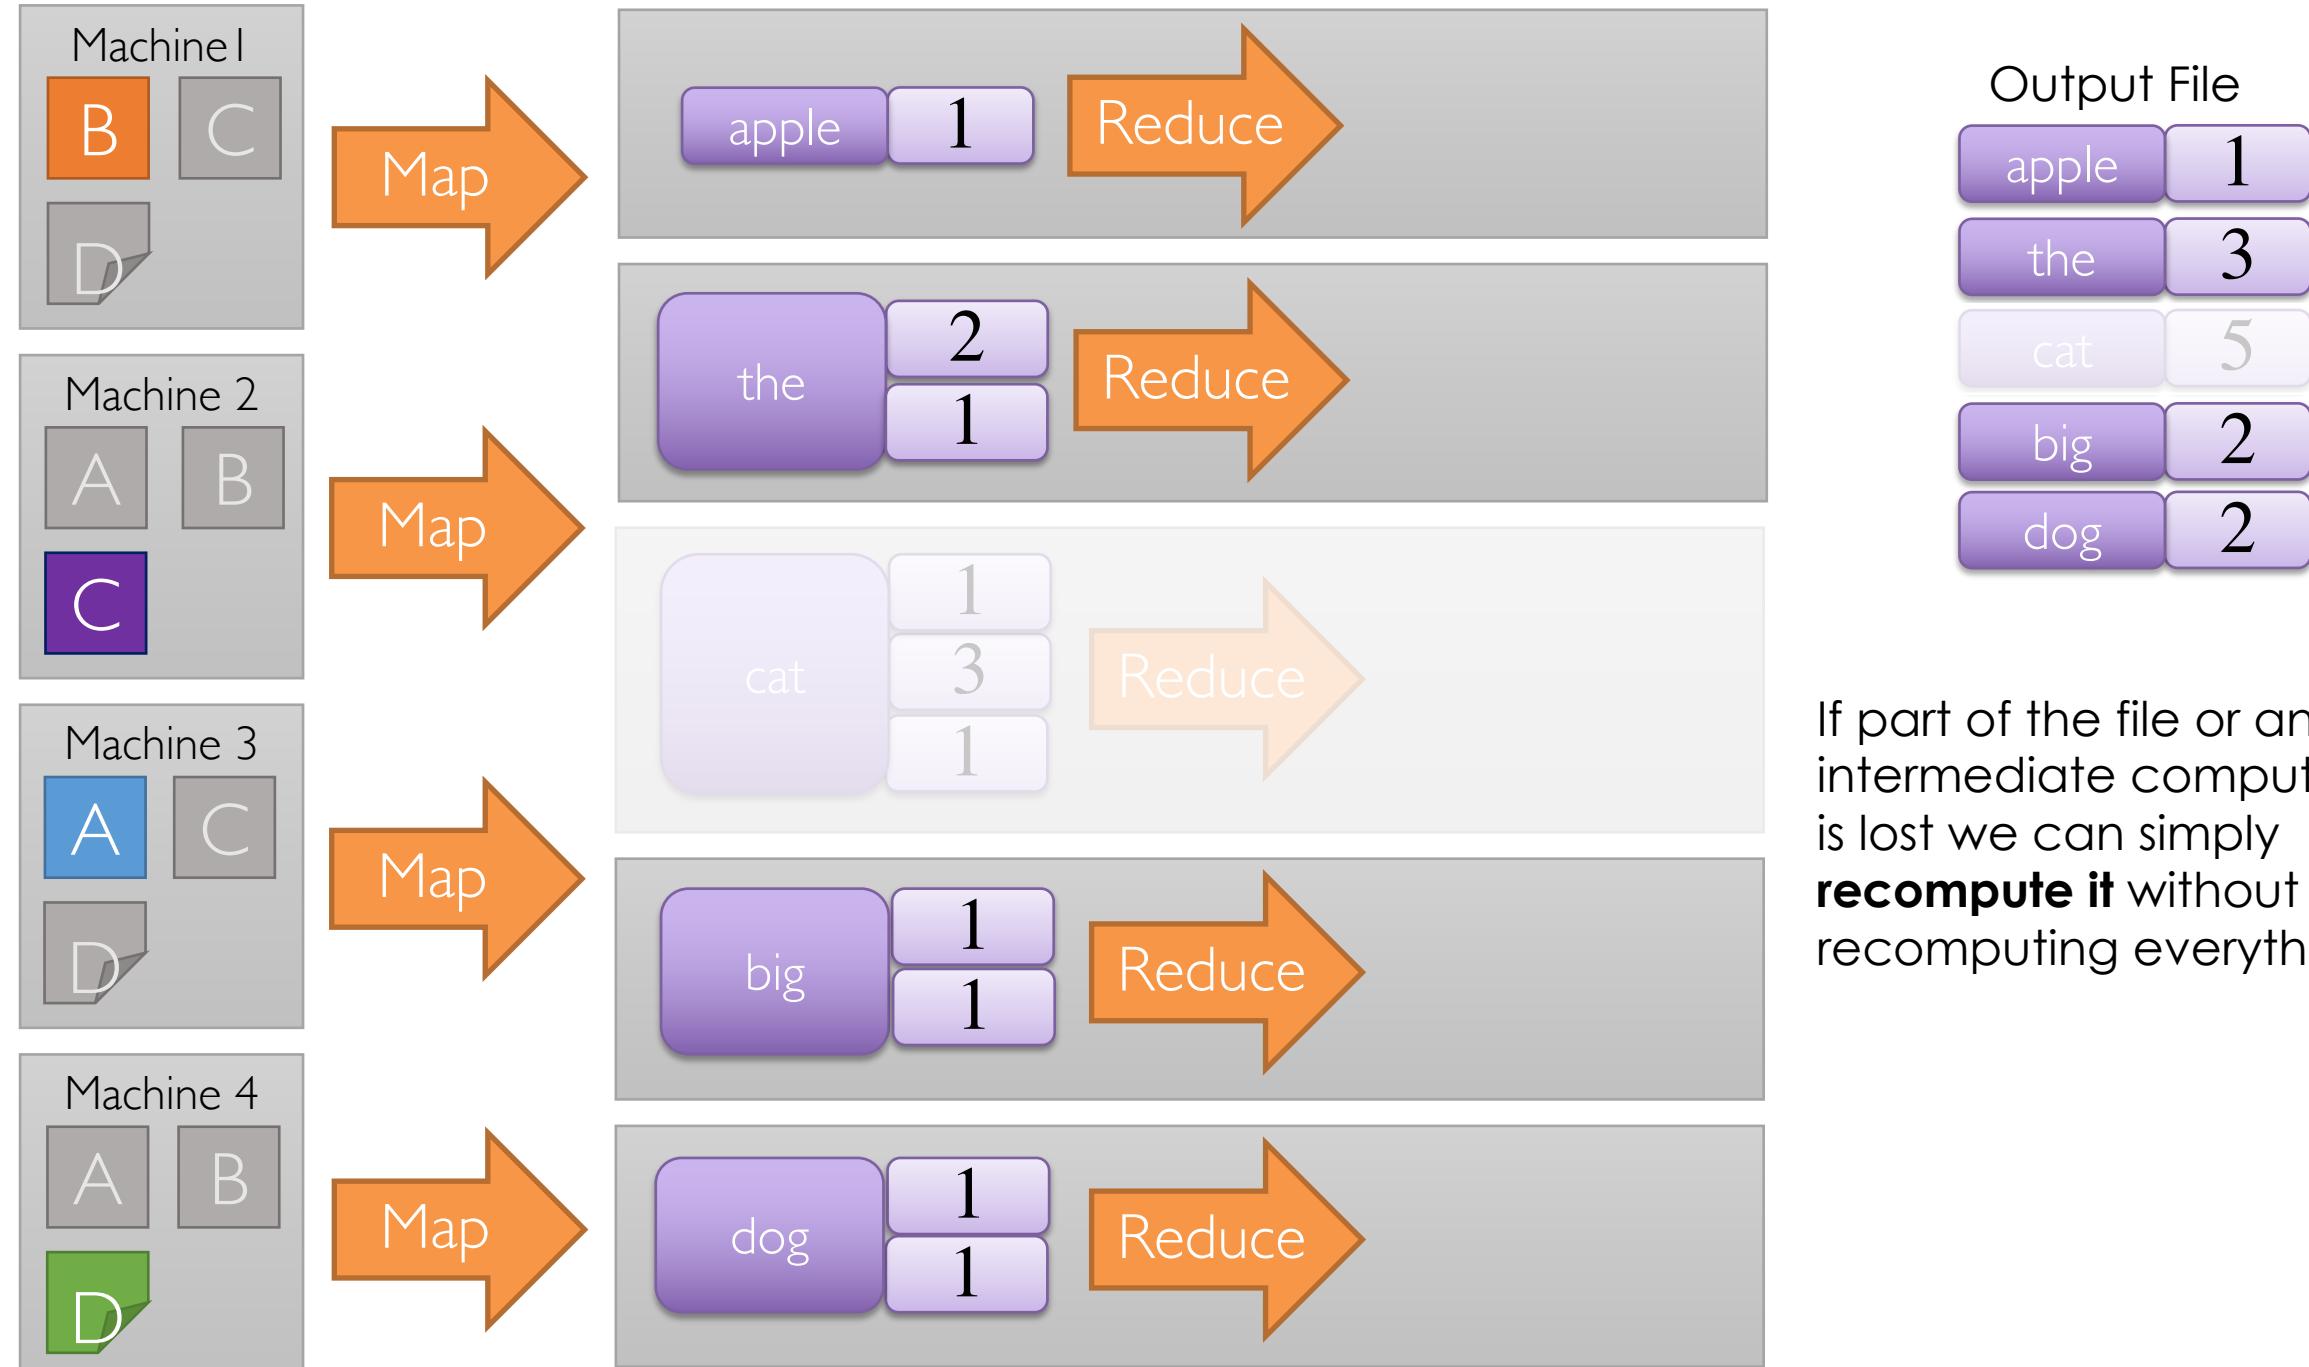


Reduce function can be run on many machines ...

Run in Parallel





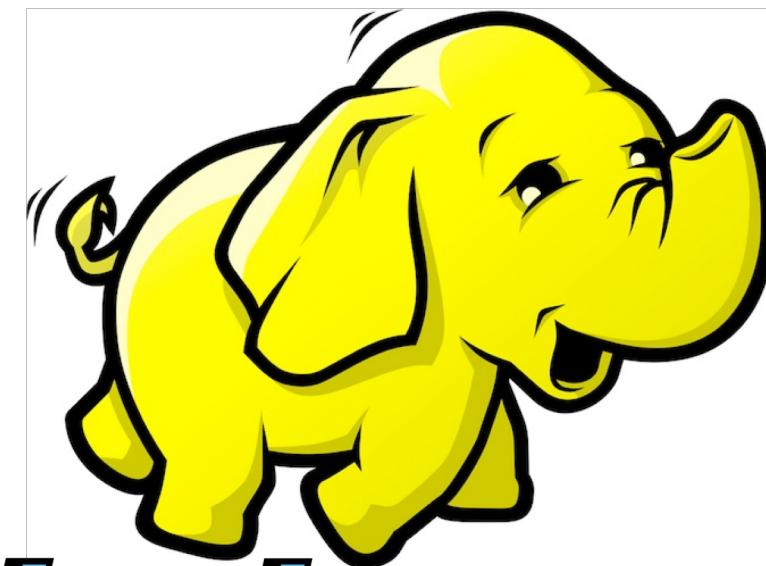


If part of the file or any intermediate computation is lost we can simply **recompute it** without recomputing everything.

Map Reduce Technologies

Hadoop

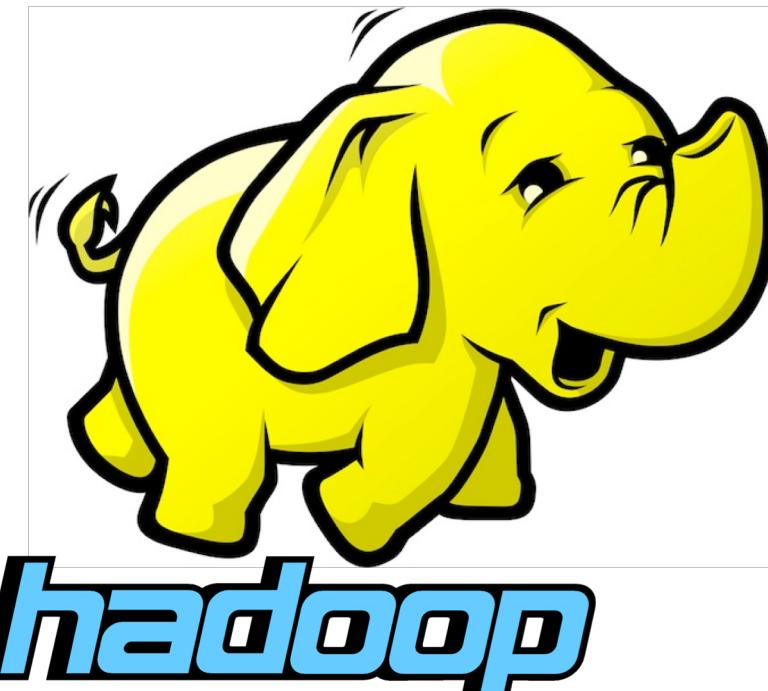
- First open-source map-reduce software
 - Managed by Apache foundation
- Based on Google's
 - Google File System
 - MapReduce
- Companies formed around Hadoop:
 - Cloudera
 - Hortonworks
 - MapR



hadoop

Hadoop

- Very active open source ecosystem
- Several key technologies
 - **HDFS**: Hadoop File System
 - **MapReduce**: map-reduce compute framework
 - **YARN**: Yet another resource negotiator
 - **Hive**: SQL queries over MapReduce
 - ...
- Downside: Tedious to use!
 - Joey: Word count example from before is 100s of lines of Java code.





In-Memory Dataflow System

Developed at the UC Berkeley AMP Lab

M. Zaharia, M. Chowdhury, M. J. Franklin, S. Shenker, and I. Stoica. *Spark: cluster computing with working sets*. HotCloud'10

M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauley, M.J. Franklin, S. Shenker, I. Stoica. *Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing*, NSDI 2012

What Is Spark

- Parallel execution engine for big data processing
- **General**: efficient support for multiple workloads
- **Easy** to use: 2-5x less code than Hadoop MR
 - High level API's in Python, Java, and Scala
- **Fast**: up to 100x faster than Hadoop MR
 - Can exploit in-memory when available
 - Low overhead scheduling, optimized engine

Spark Programming Abstraction

- Write programs *in terms of transformations on distributed datasets*
- Resilient Distributed Datasets ([RDDs](#))
 - Distributed [collections of objects](#) that can be stored [in memory](#) or [on disk](#)
 - Built via parallel transformations (map, filter, ...)
 - Automatically rebuilt on failure

RDD: Resilient Distributed Datasets

- Collections of objects partitioned & distributed across a cluster
 - Stored in RAM or on Disk
 - Resilient to failures
- Operations
 - Transformations
 - Actions

Operations on RDDs

- Transformations $f(\text{RDD}) \Rightarrow \text{RDD}$
 - Lazy (not computed immediately)
 - E.g., “map”, “filter”, “groupBy”

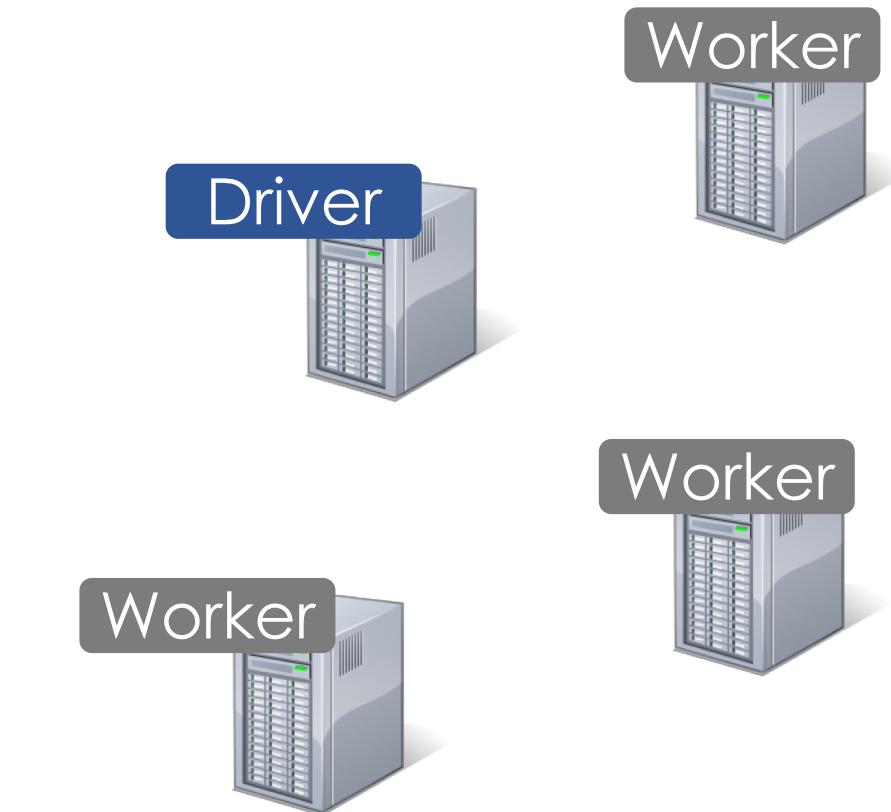
- Actions:
 - Triggers computation
 - E.g. “count”, “collect”, “saveAsTextFile”

Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

Example: Log Mining

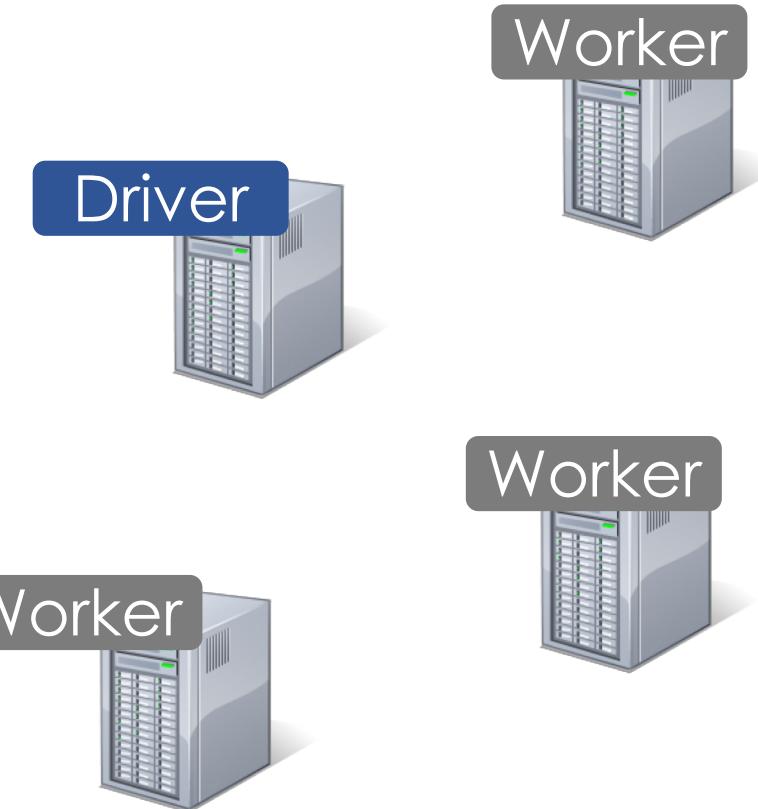
Load error messages from a log into memory,
then interactively search for various patterns



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

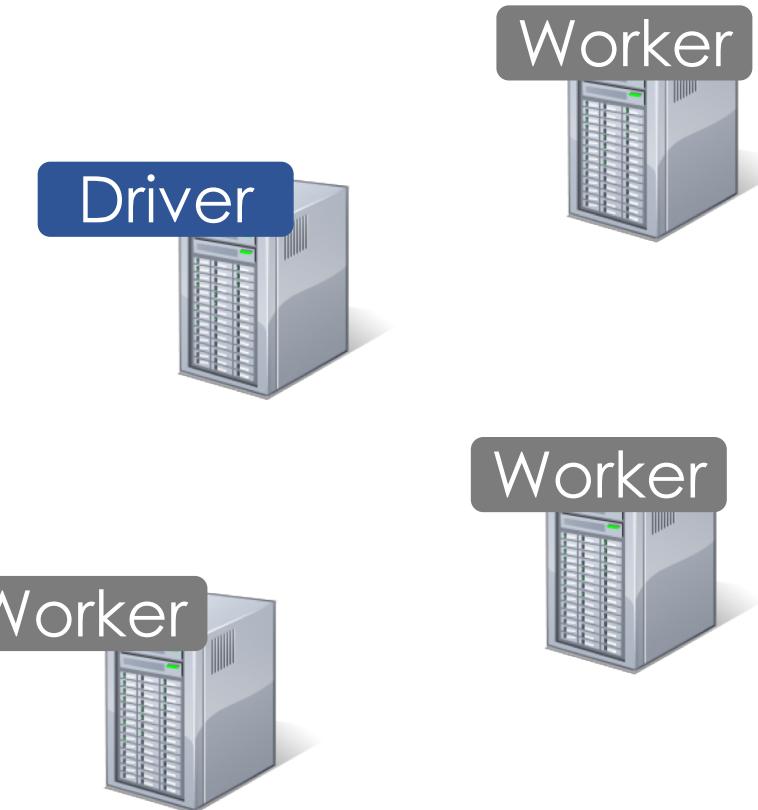
```
lines = spark.textFile("hdfs://file.txt")
```



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

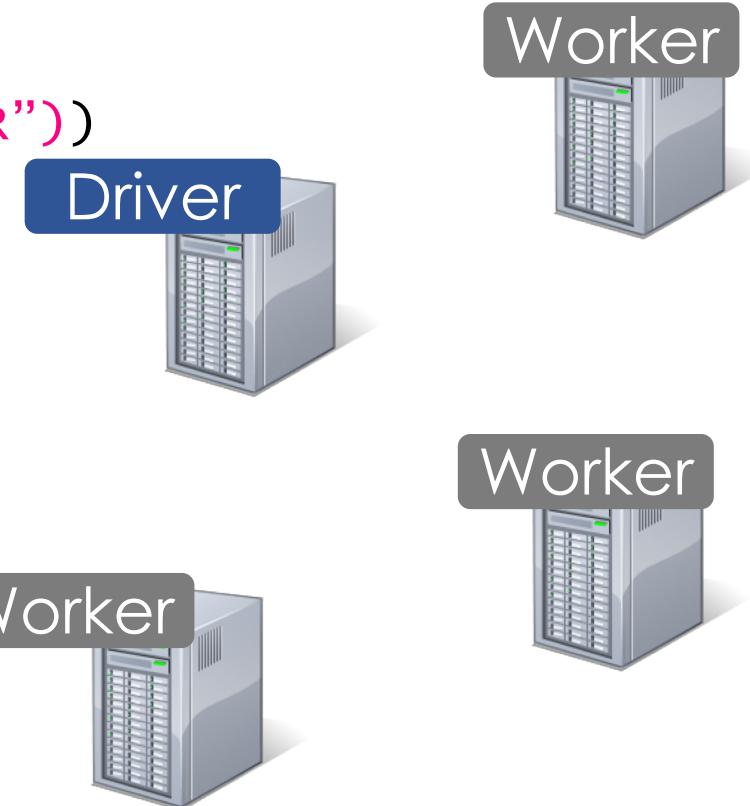
Base RDD
lines = spark.textFile("hdfs://file.txt")



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
```



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

Transformed RDD

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
```

Driver



Worker



Worker



Worker

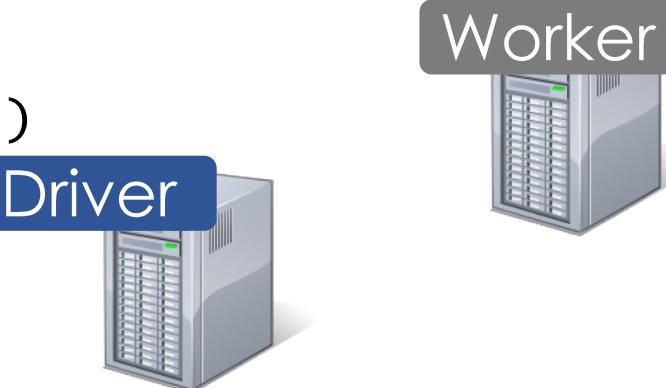


Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

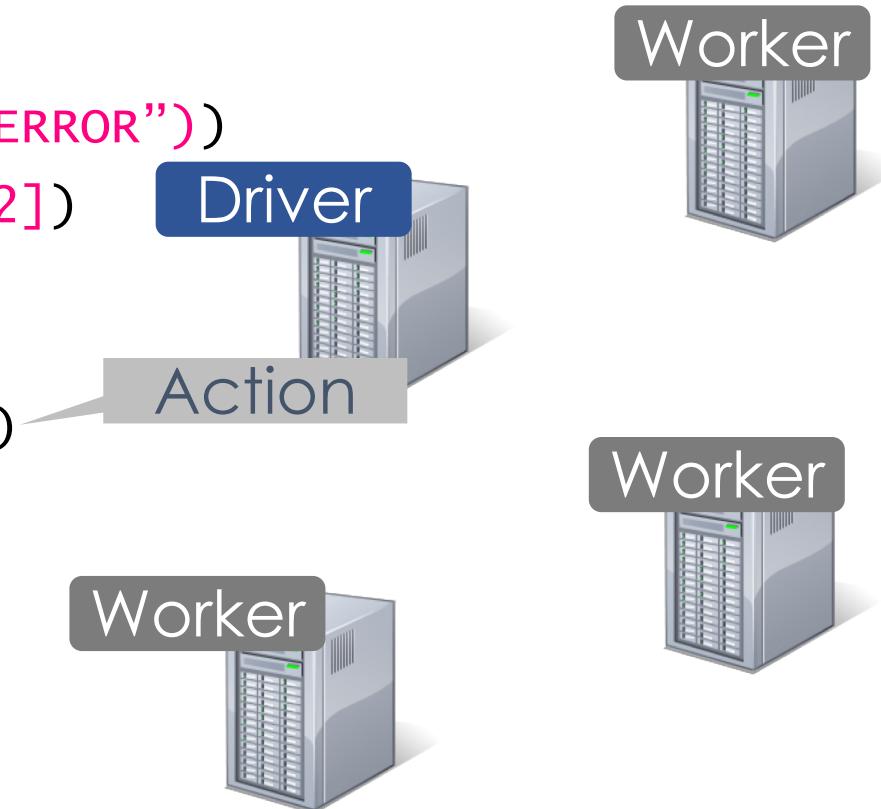
messages.filter(lambda s: "mysql" in s).count()
```



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
messages.filter(lambda s: "mysql" in s).count()
```

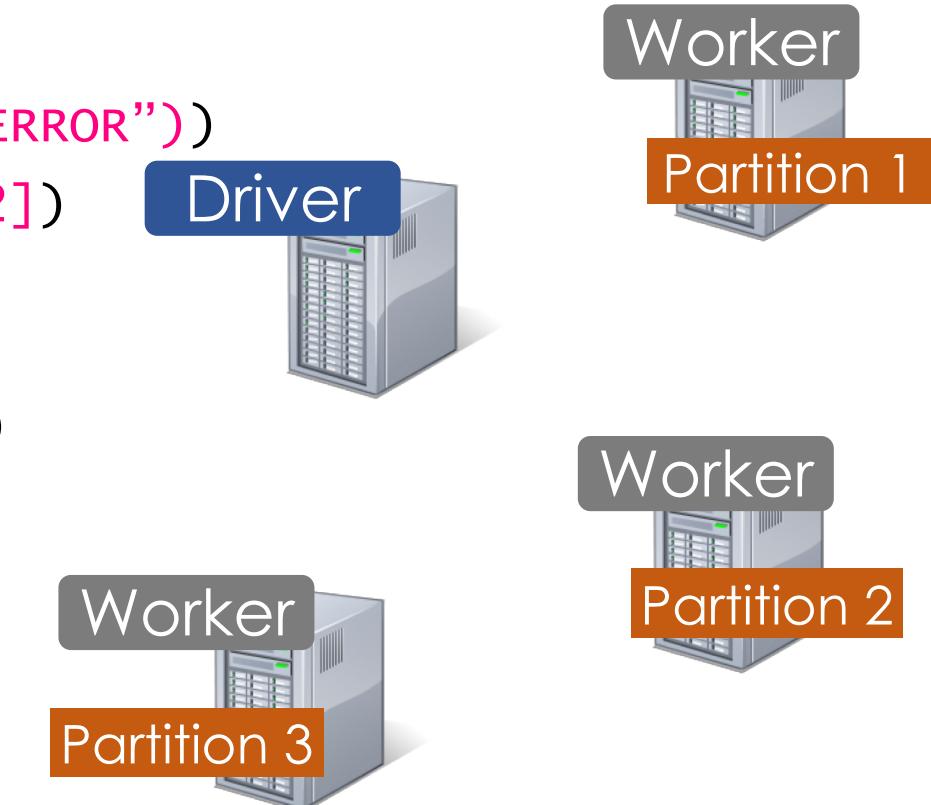


Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
```

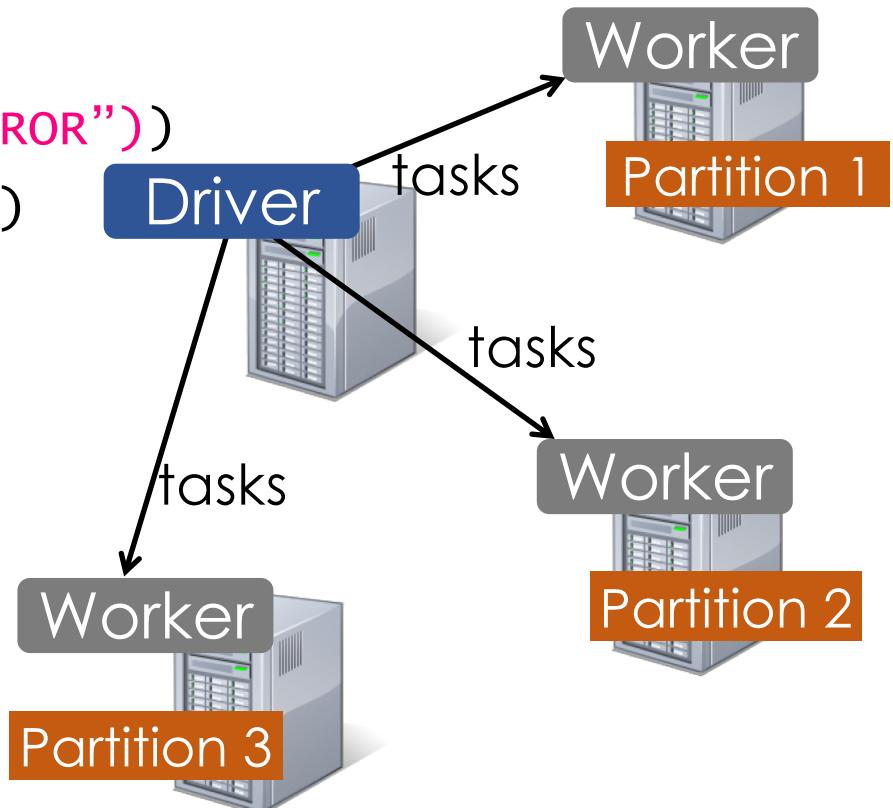


Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
```

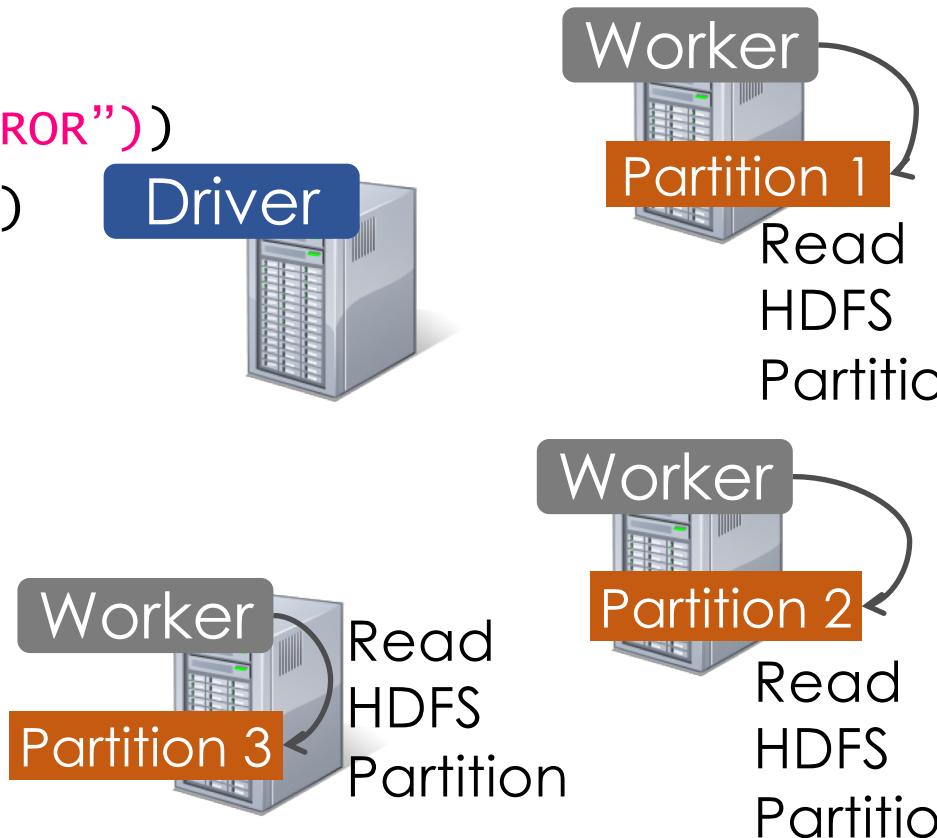


Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

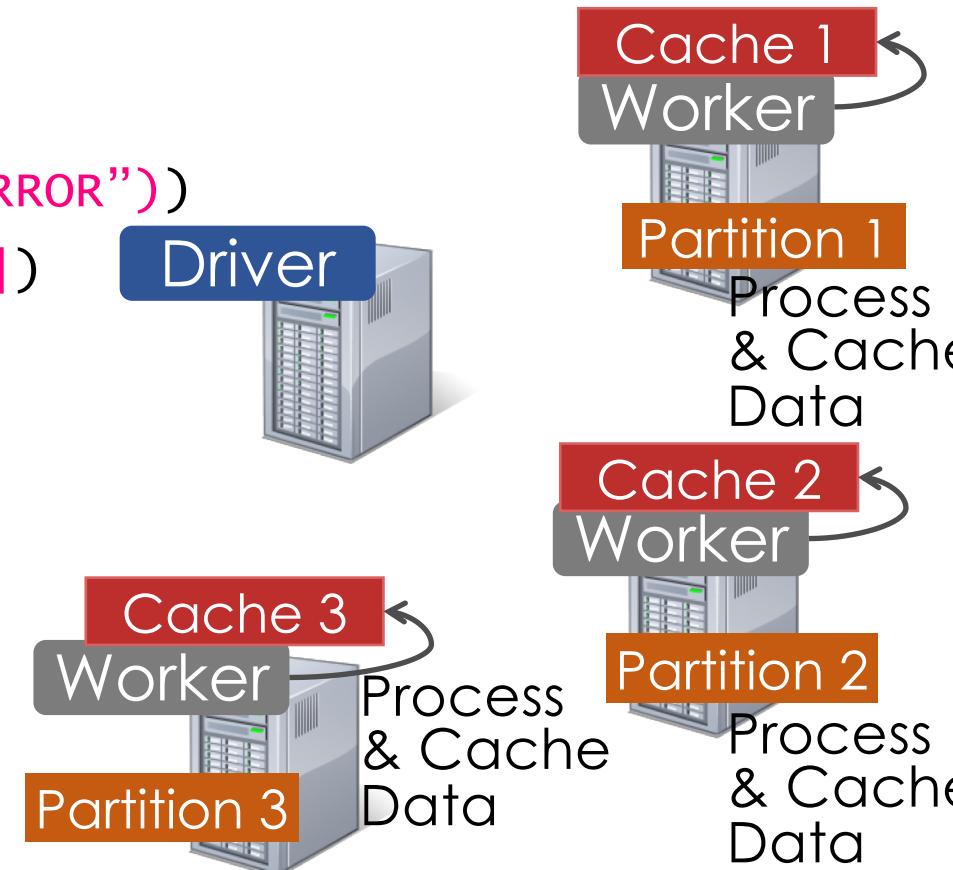
messages.filter(lambda s: "mysql" in s).count()
```



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()
messages.filter(lambda s: "mysql" in s).count()
```

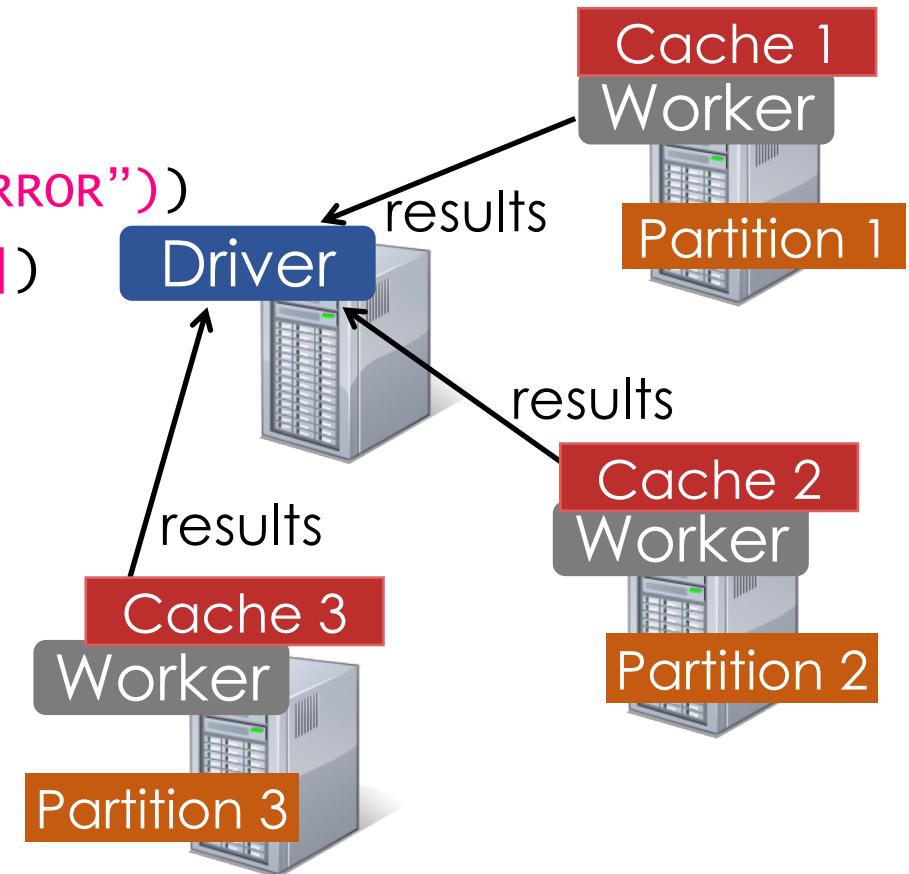


Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
```

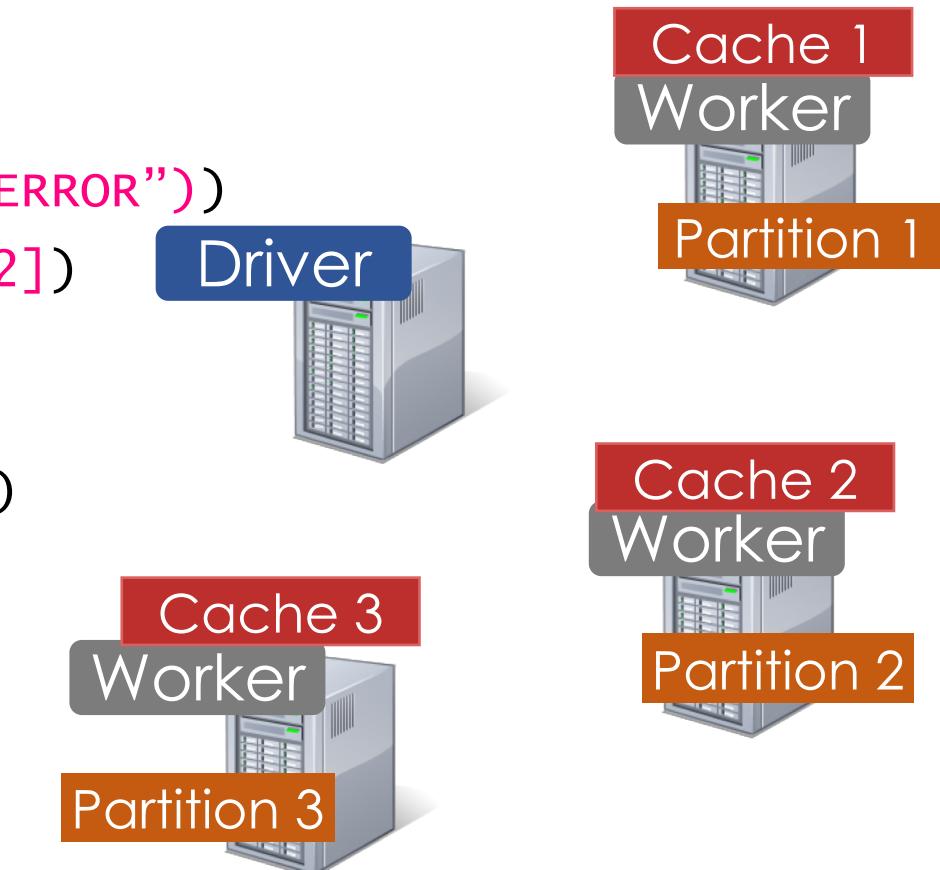


Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

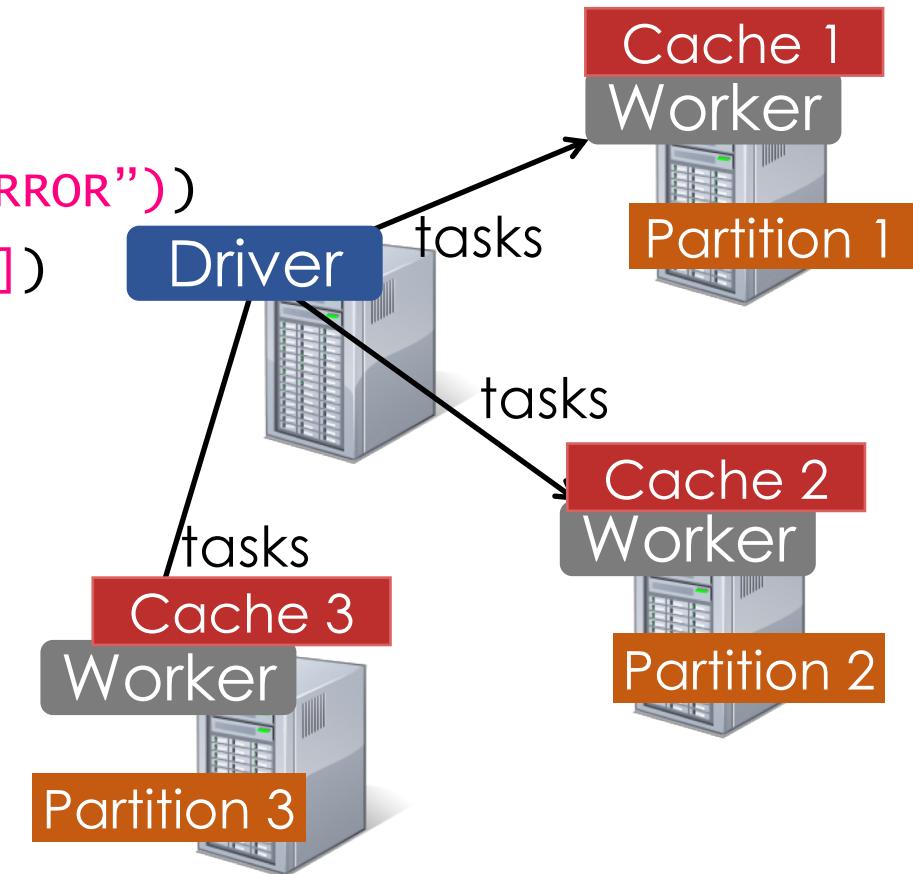


Example: Log Mining

Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

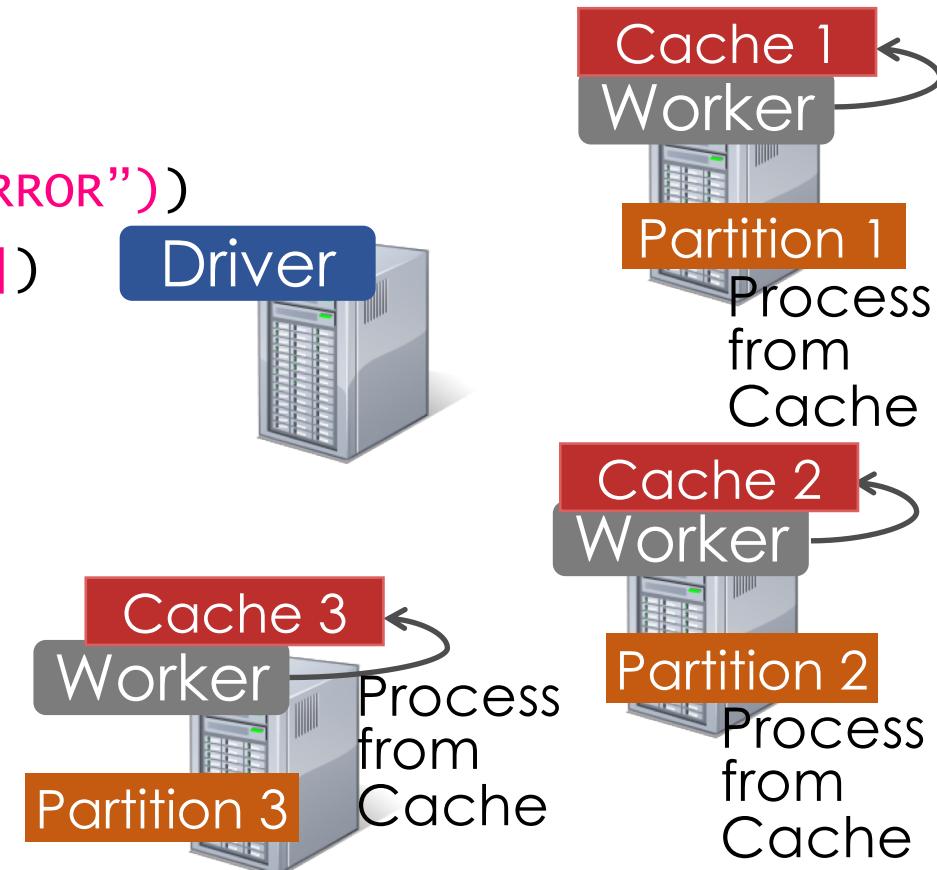


Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

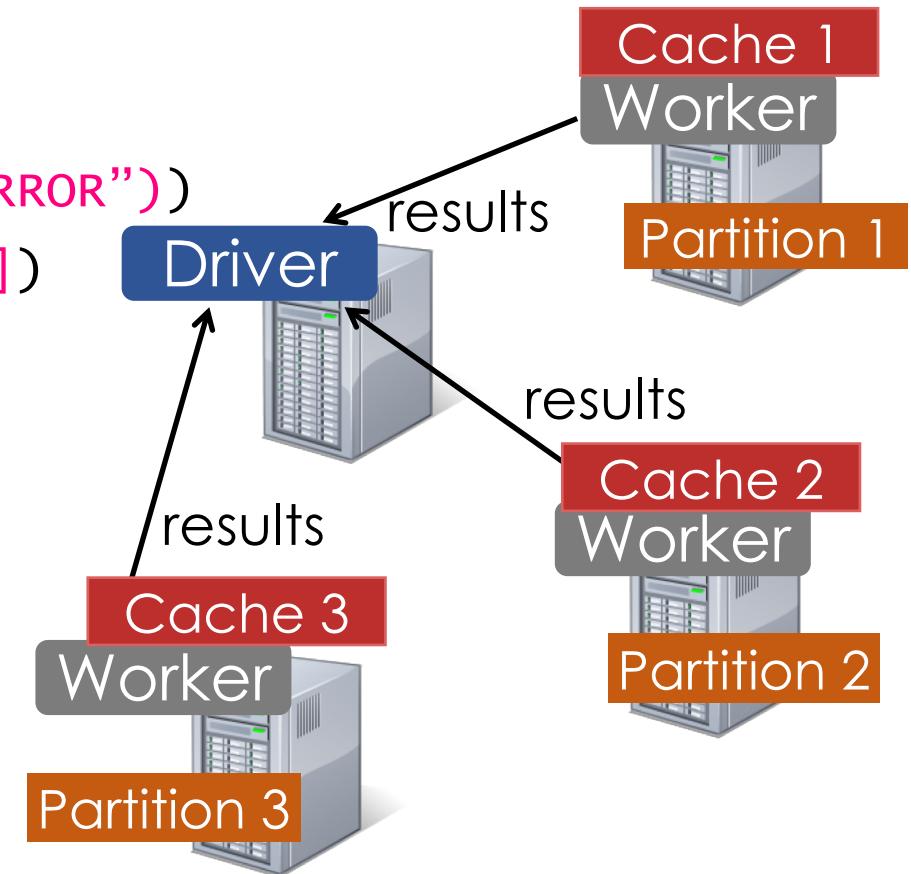


Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```



Example: Log Mining

Load error messages from a log into memory,
then interactively search for various patterns

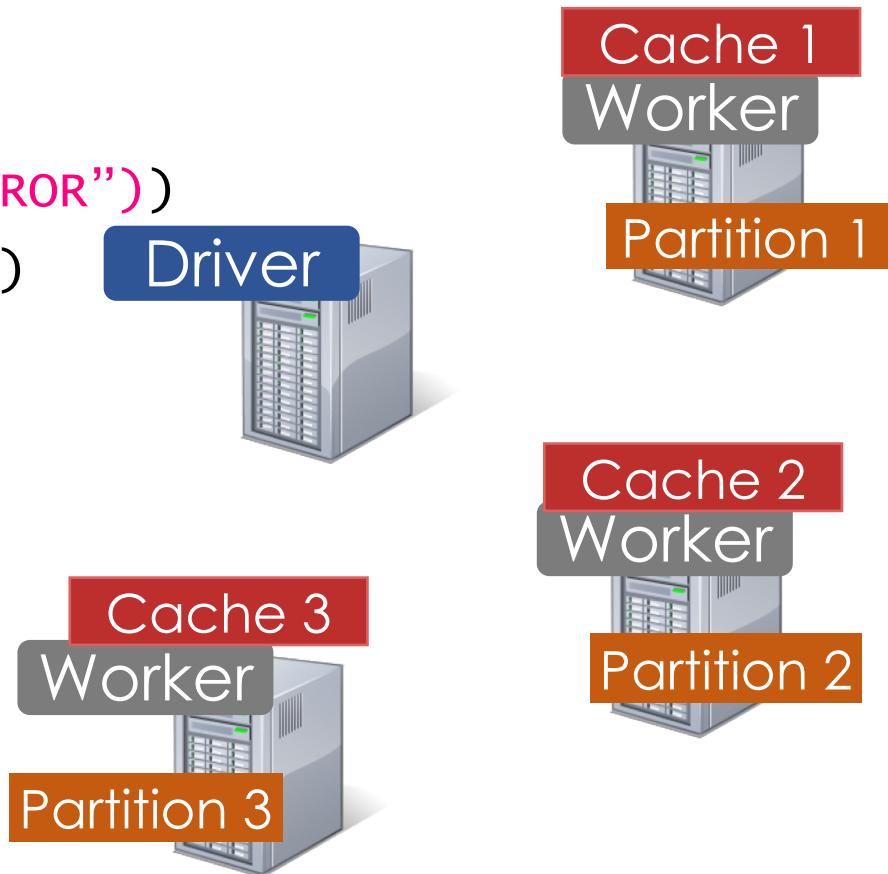
```
lines = spark.textFile("hdfs://file.txt")
errors = lines.filter(lambda s: s.startswith("ERROR"))
messages = errors.map(lambda s: s.split("\t")[2])
messages.cache()

messages.filter(lambda s: "mysql" in s).count()
messages.filter(lambda s: "php" in s).count()
```

Cache your data → Faster Results

Full-text search of Wikipedia

- 60GB on 20 EC2 machines
- 0.5 sec from mem vs. 20s for on-disk



Example: Counting Words

```
moby_dick = spark.textFile("hdfs://books/mobydick.txt")
lines = moby_dick.flatMap(lambda line: line.split(" "))
counts = lines.map(lambda word: (word, 1))
    .reduceByKey(lambda word: (word, 1))

counts.toDF().toPandas()
```

Much simpler than
Hadoop map
reduce code!

- The **flatMap** and **map** calls produce transformed RDDs.
- Computation is lazy! Nothing happens until we get to our action.
- The **reduceByKey** calls kicks off the actual computing.

Abstraction: Dataflow Operators

map

filter

groupBy

sort

union

join

leftOuterJoin

rightOuterJoin

reduce

count

fold

reduceByKey

groupByKey

cogroup

cross

zip

sample

take

first

partitionBy

mapWith

pipe

save

...

Abstraction: Dataflow Operators

map

filter

groupBy

sort

union

join

leftOuterJoin

rightOuterJoin

reduce

count

fold

reduceByKey

groupByKey

cogroup

cross

zip

sample

take

first

partitionBy

mapwith

pipe

save

...

Spark Demo

Summary (1/2)

- ETL is used to bring data from operational data stores into a data warehouse.
 - Many ways to organize tabular data warehouse, e.g. star and snowflake schemas.
- Online Analytics Processing (OLAP) techniques let us analyze data in data warehouse.
 - Examples: Pivot table, CUBE, slice, dice, rollup, drill down.
- Unstructured data is hard to store in a tabular format in a way that is amenable to standard techniques, e.g. finding pictures of cats.
 - Resulting new paradigm: The Data Lake.

Summary (2/2)

- Data Lake is enabled by two key ideas:
 - Distributed file storage.
 - Distributed computation.
- Distributed file storage involves replication of data.
 - Better speed and reliability, but more costly.
- Distributed computation made easier by map reduce.
 - Hadoop: Open-source implementation of distributed file storage and computation.
 - Spark: Typically faster and easier to use than Hadoop.